

Image Domain Conversion

Upsampling Alternatives and Applications

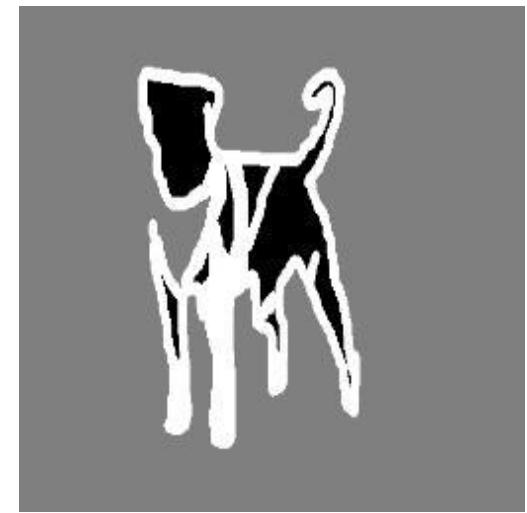
Quang-Vinh Dinh
Ph.D. in Computer Science

Outline

- Segmentation (Unet Using Interpolation)
- Alternatives to Increase Feature Resolution
- Colorization
- Super-resolution
- Denoising Survey
- Super-resolution Survey

Discussion – Segmentation Problem

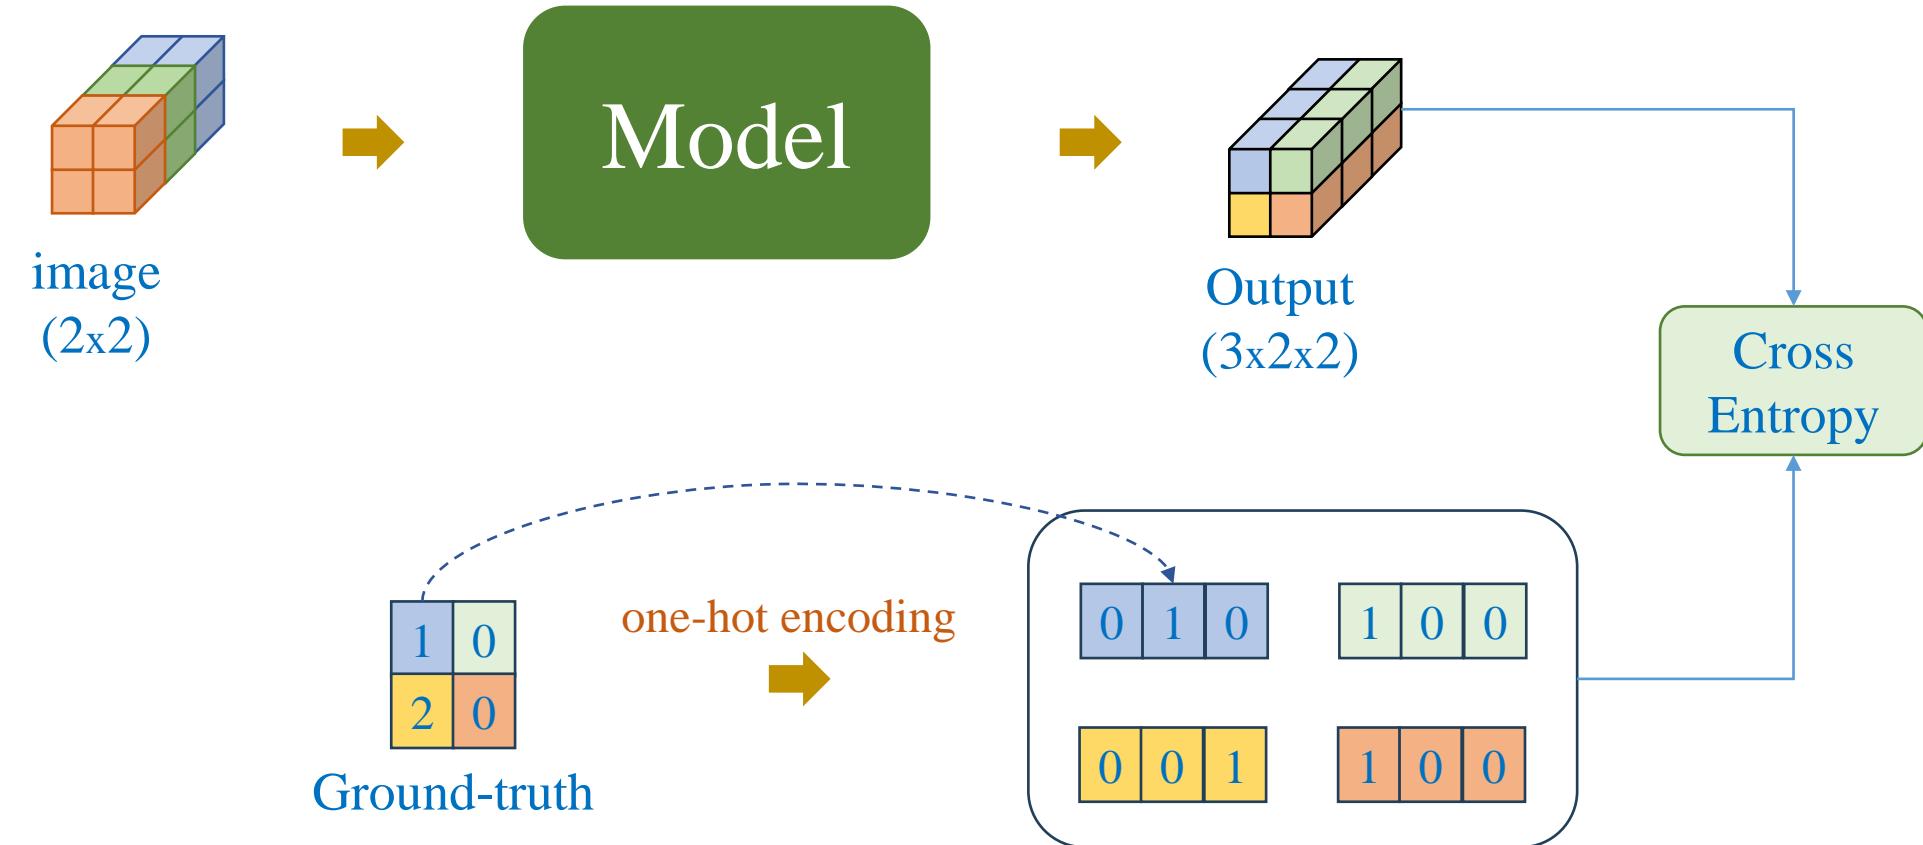
Input/Output + Loss Function



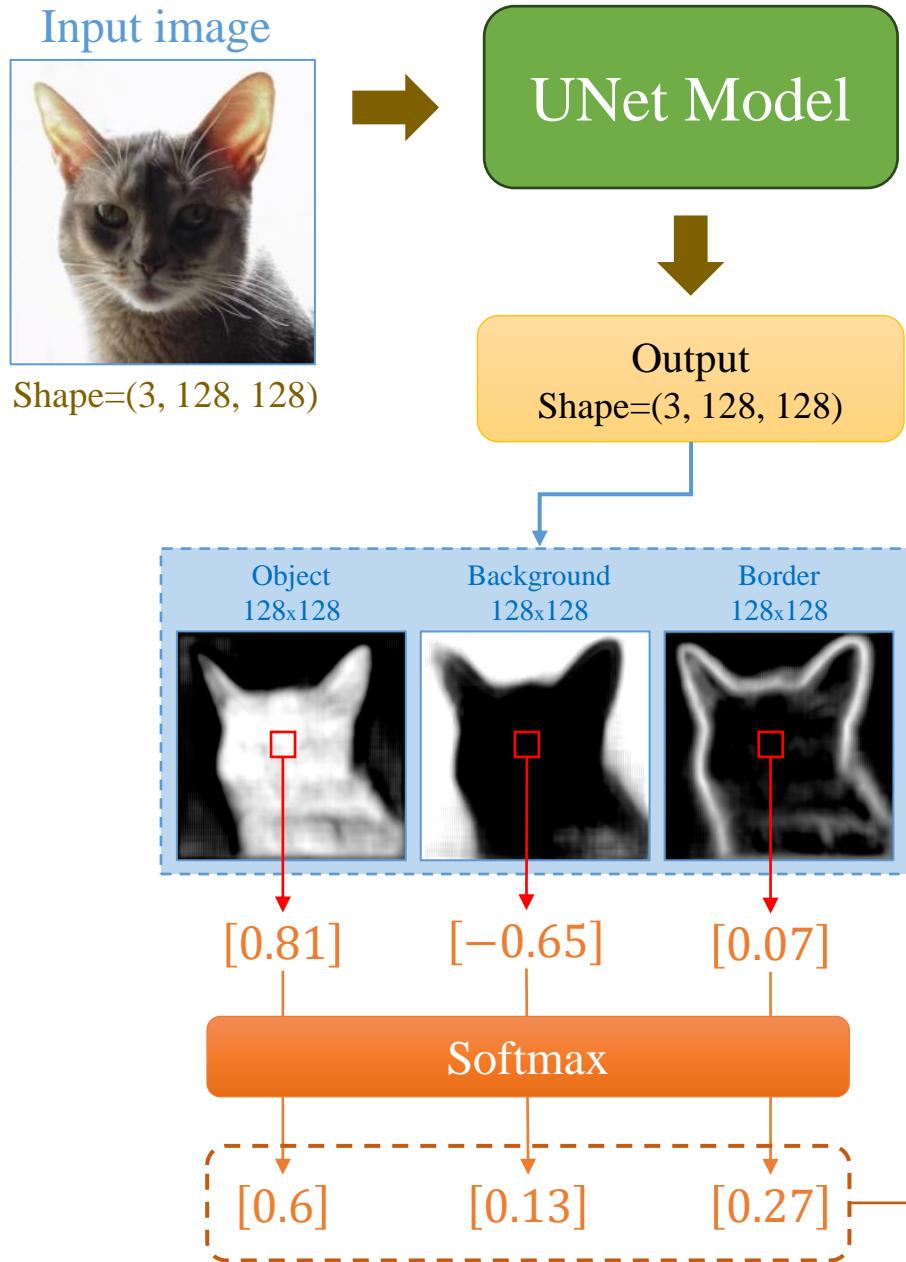
Inference



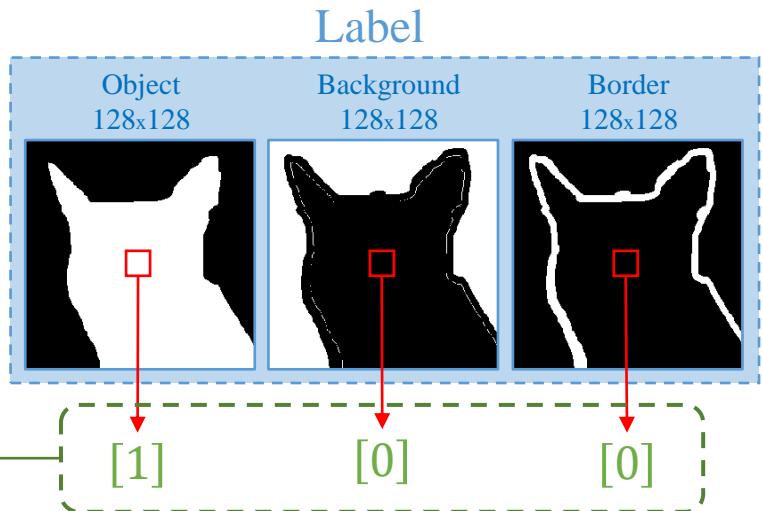
Training



Loss Function



$$CE = - \sum_i y_i \log(\hat{y}_i)$$



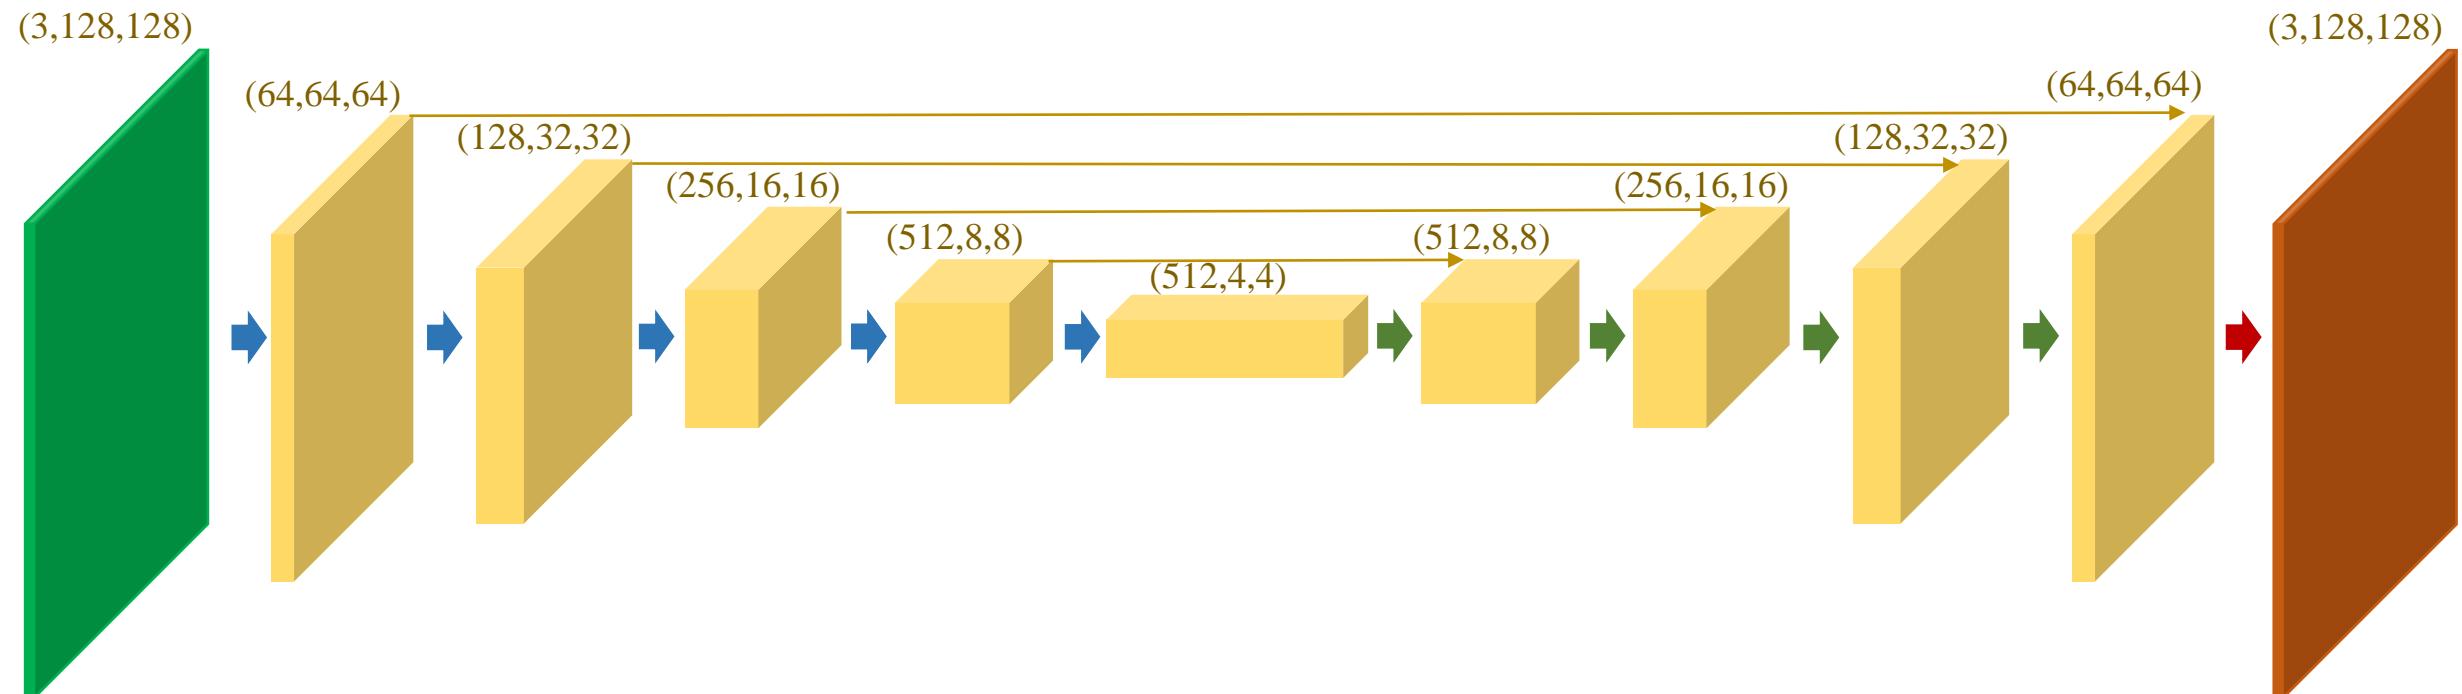
$$CE \left(\begin{bmatrix} 0.60 \\ 0.13 \\ 0.27 \end{bmatrix}, \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} \right) = 0.5$$

$$\text{loss} = \frac{1}{H * W} \sum_{j=1}^H \sum_{k=1}^W CE(\hat{y}_{j,k}, y_{j,k})$$

Segmentation (6)

Using skip connections

Code Reading



(3x3) Convolution
padding = 'same'
stride = 1 + ReLU

+ Batch
Norm

+ (2x2) max
pooling

Feature Map
Upsampling

+ (3x3) Convolution
padding = 'same'
stride = 1 + ReLU

+ Batch
Norm

Feature Map
Upsampling

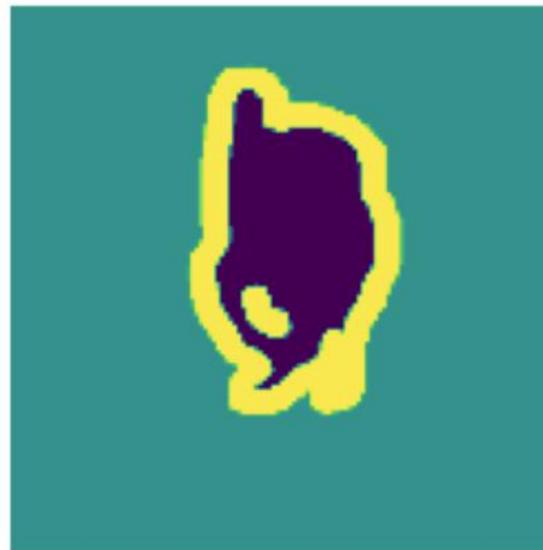
+ (3x3) Convolution
padding = 'same'
softmax

Experimental Results

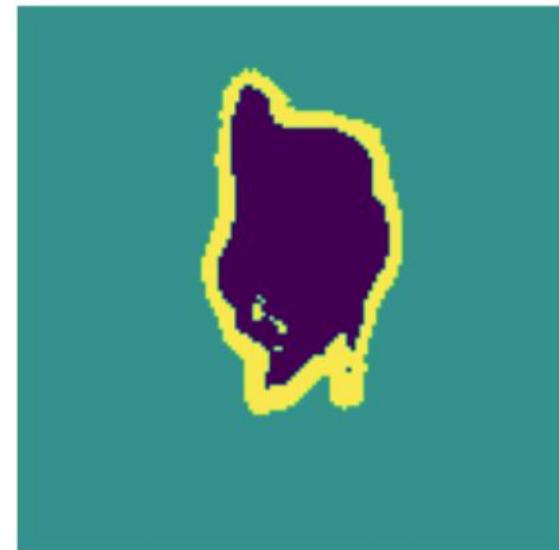
With Skip
Connections



Input Image

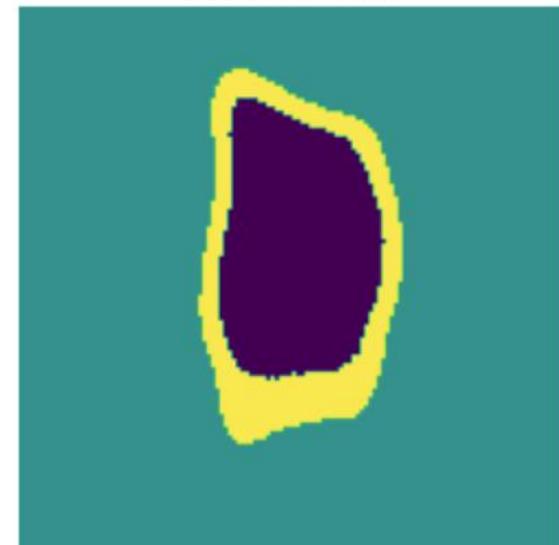
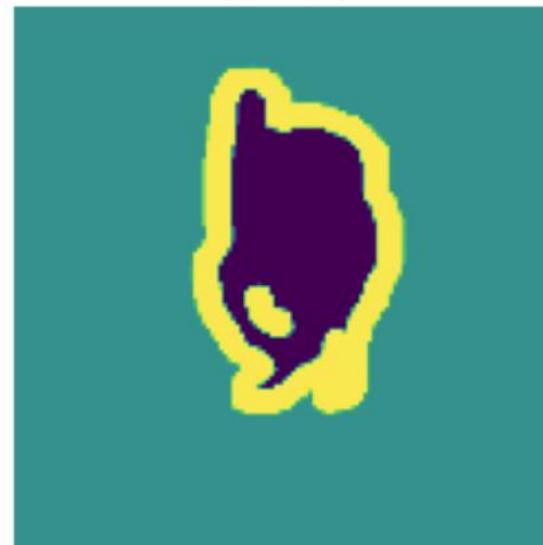


True Mask



Predicted Mask

Without Skip
Connections



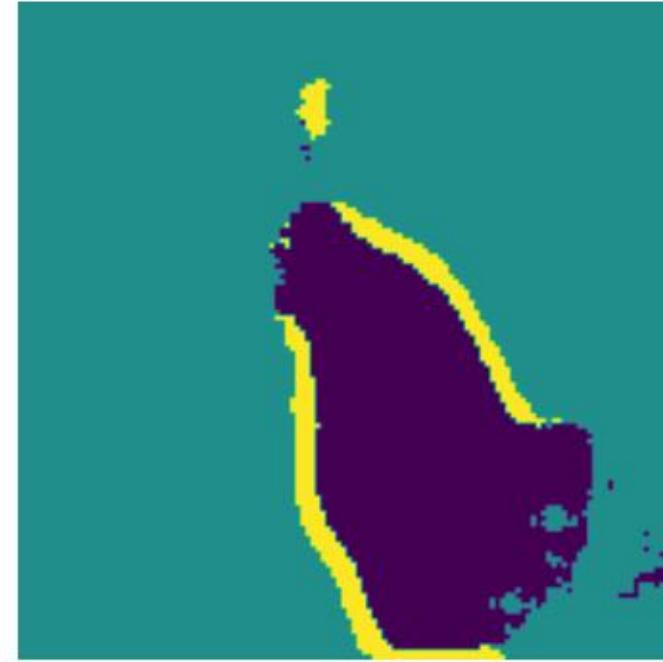
Input Image



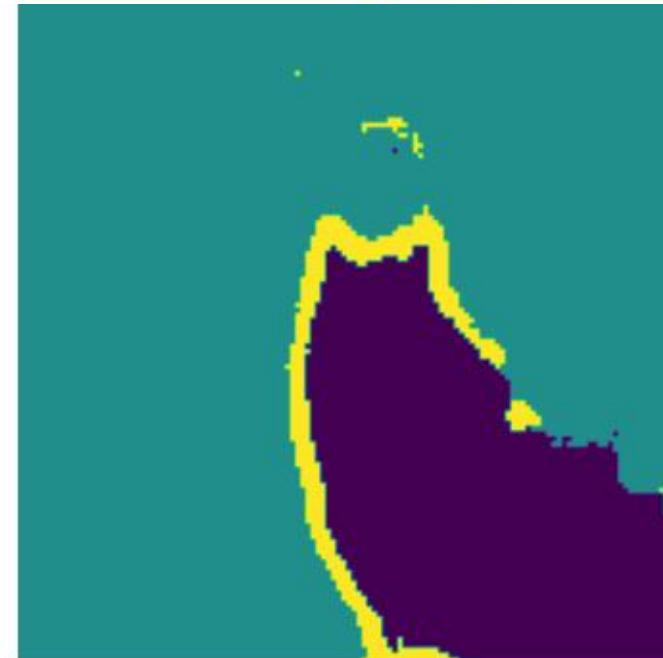
True Mask



Predicted Mask



Without
Skip
Connections



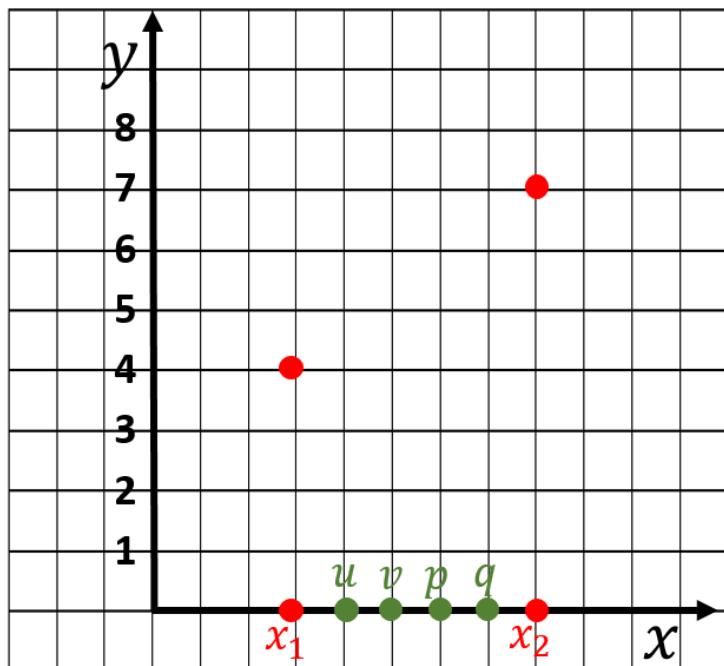
With Skip
Connections

Outline

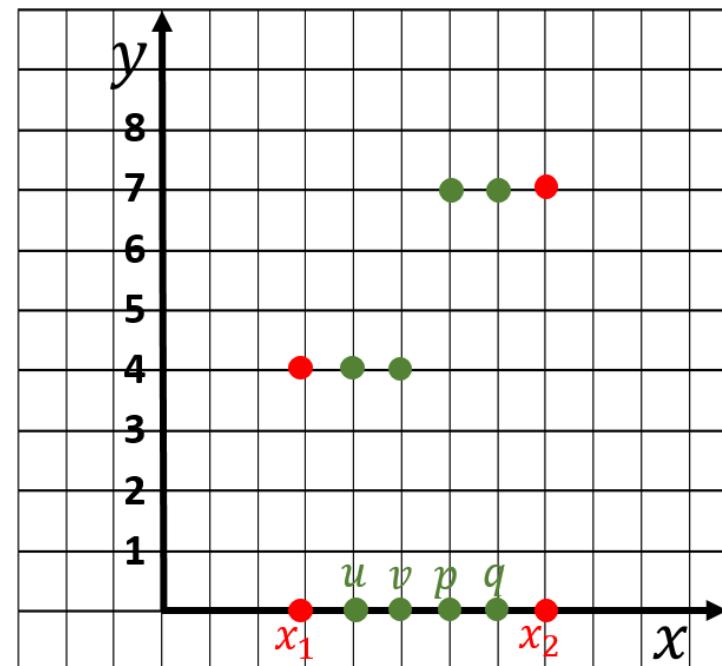
- Segmentation (Unet Using Interpolation)
- Alternatives to Increase Feature Resolution
- Colorization
- Super-resolution
- Denoising Survey
- Super-resolution Survey

How to Increase Feature Map

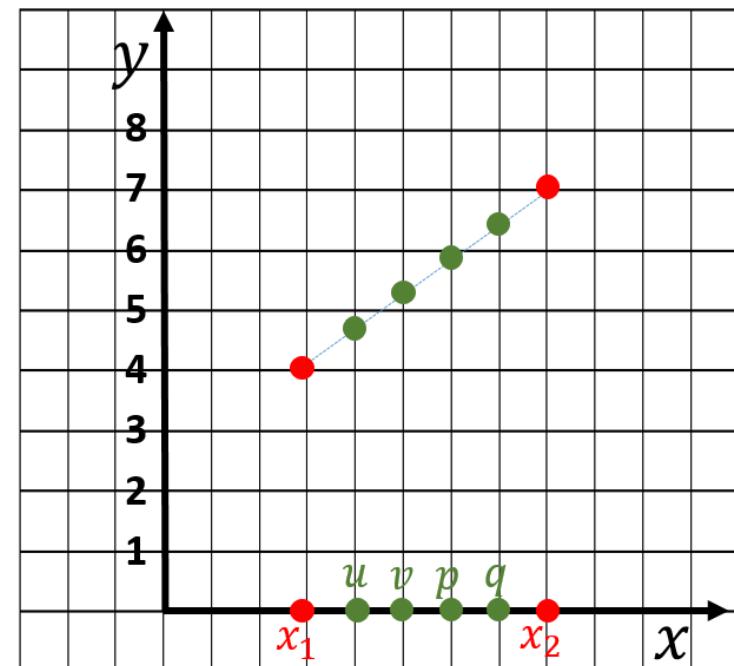
- ❖ Image upsampling
- ❖ Data interpolation



Tìm giá trị cho các vị trí u , v , p và q



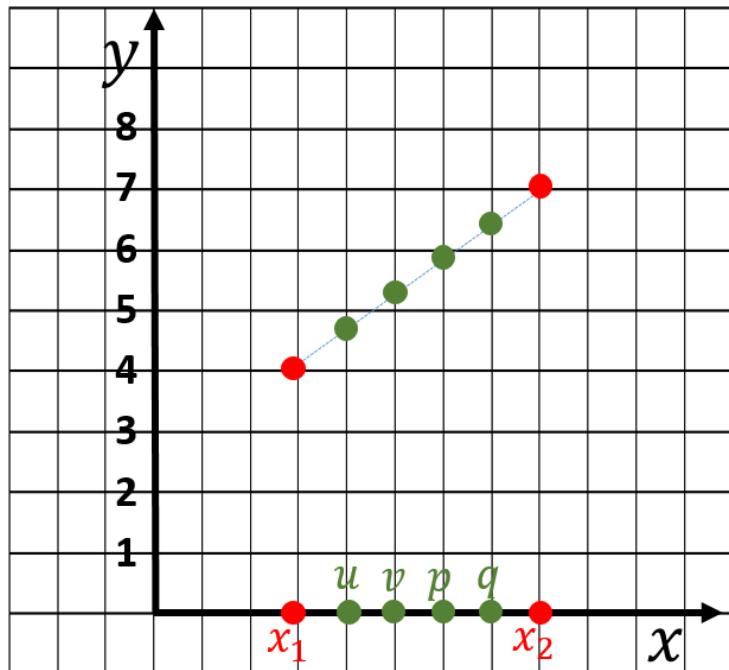
Nearest neighbor: Tính khoảng cách
đến x_1 và x_2 , và lấy giá trị của x gần hơn



Nội suy theo hàm tuyến tính

How to Increase Feature Map

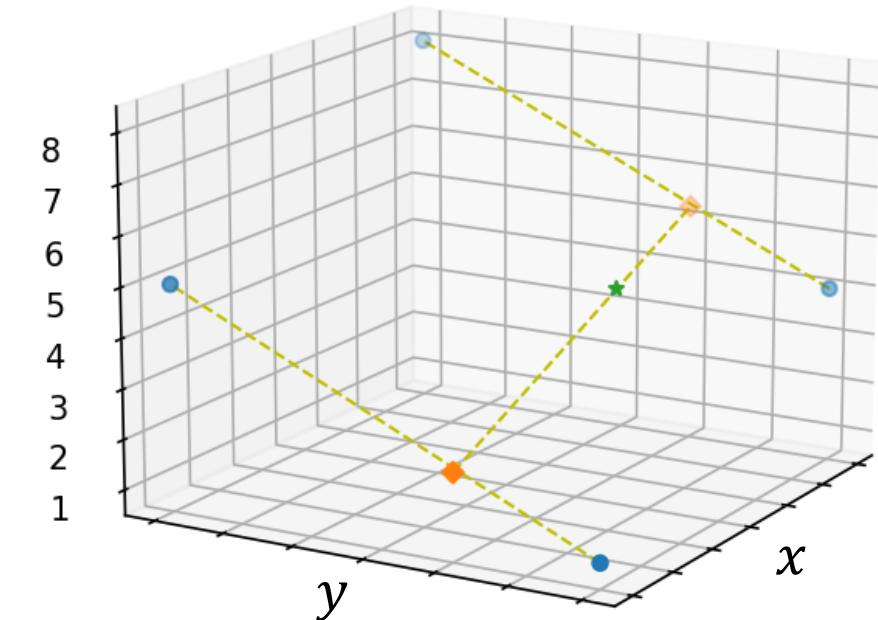
- ❖ Image upsampling
 - ❖ Data interpolation



Nội suy theo hàm tuyến tính

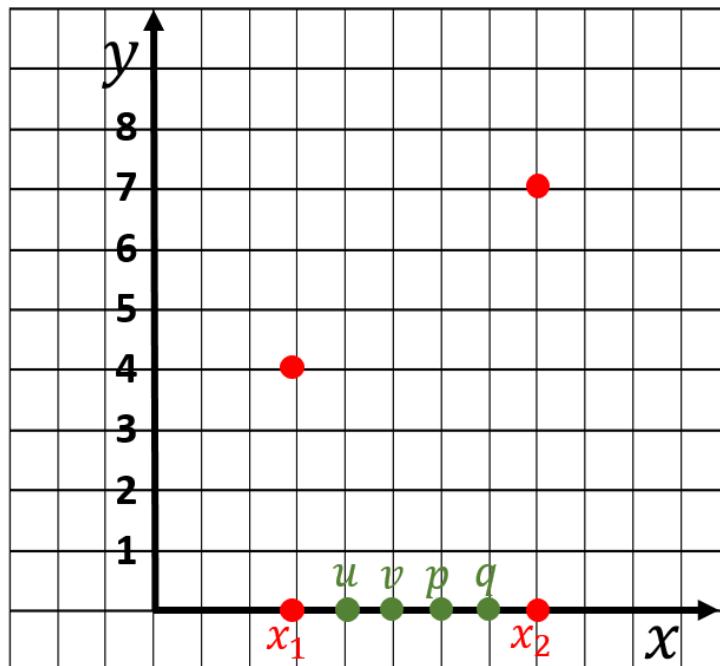
1.0			5.0
3.2	4.2		7.2
4.0			8.0

Bilinear interpolation

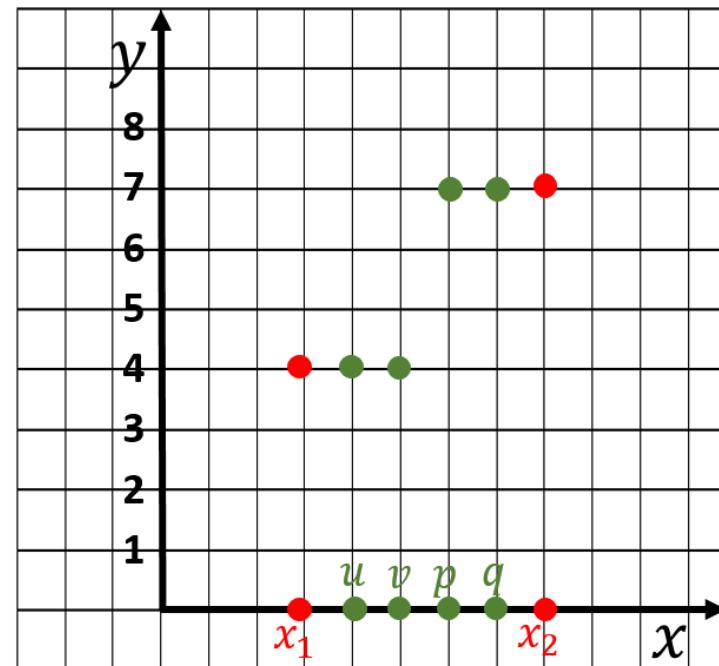


How to Increase Feature Map

❖ Nearest neighbor interpolation



Tìm giá trị cho các vị trí u, v, p và q



Nearest neighbor: Tính khoảng cách
đến x_1 và x_2 , và lấy giá trị của x gần hơn

data =

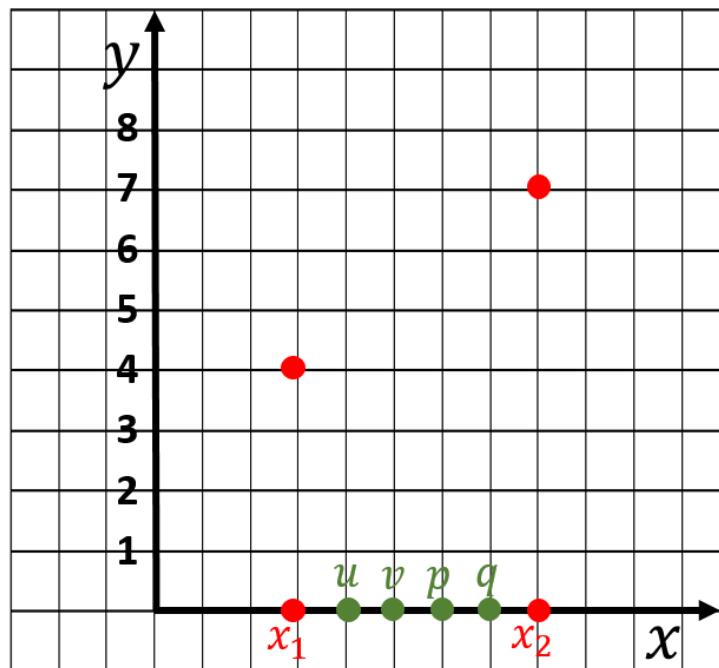
1	5
4	8

1	1	5	5
1	1	5	5
4	4	8	8
4	4	8	8

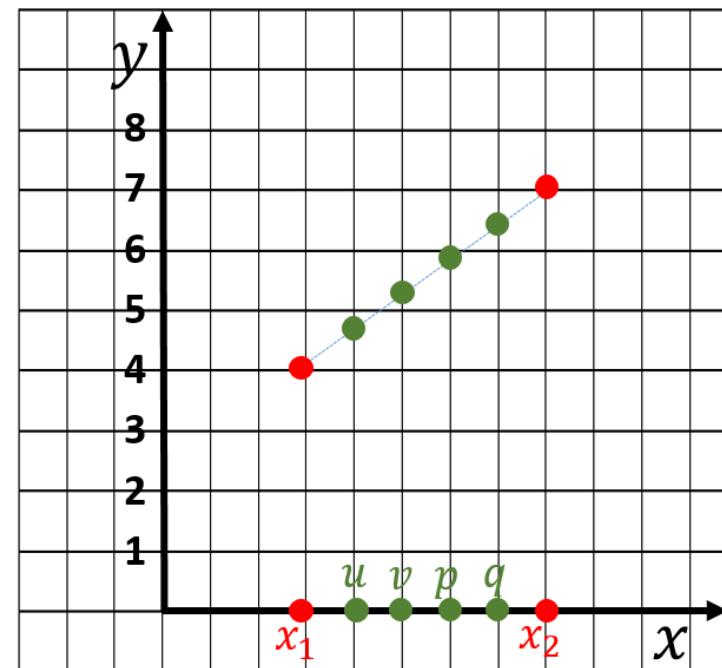
Nearest neighbor
interpolation

How to Increase Feature Map

❖ Bilinear interpolation



Tìm giá trị cho các vị trí u , v , p và q



Nội suy theo hàm tuyến tính

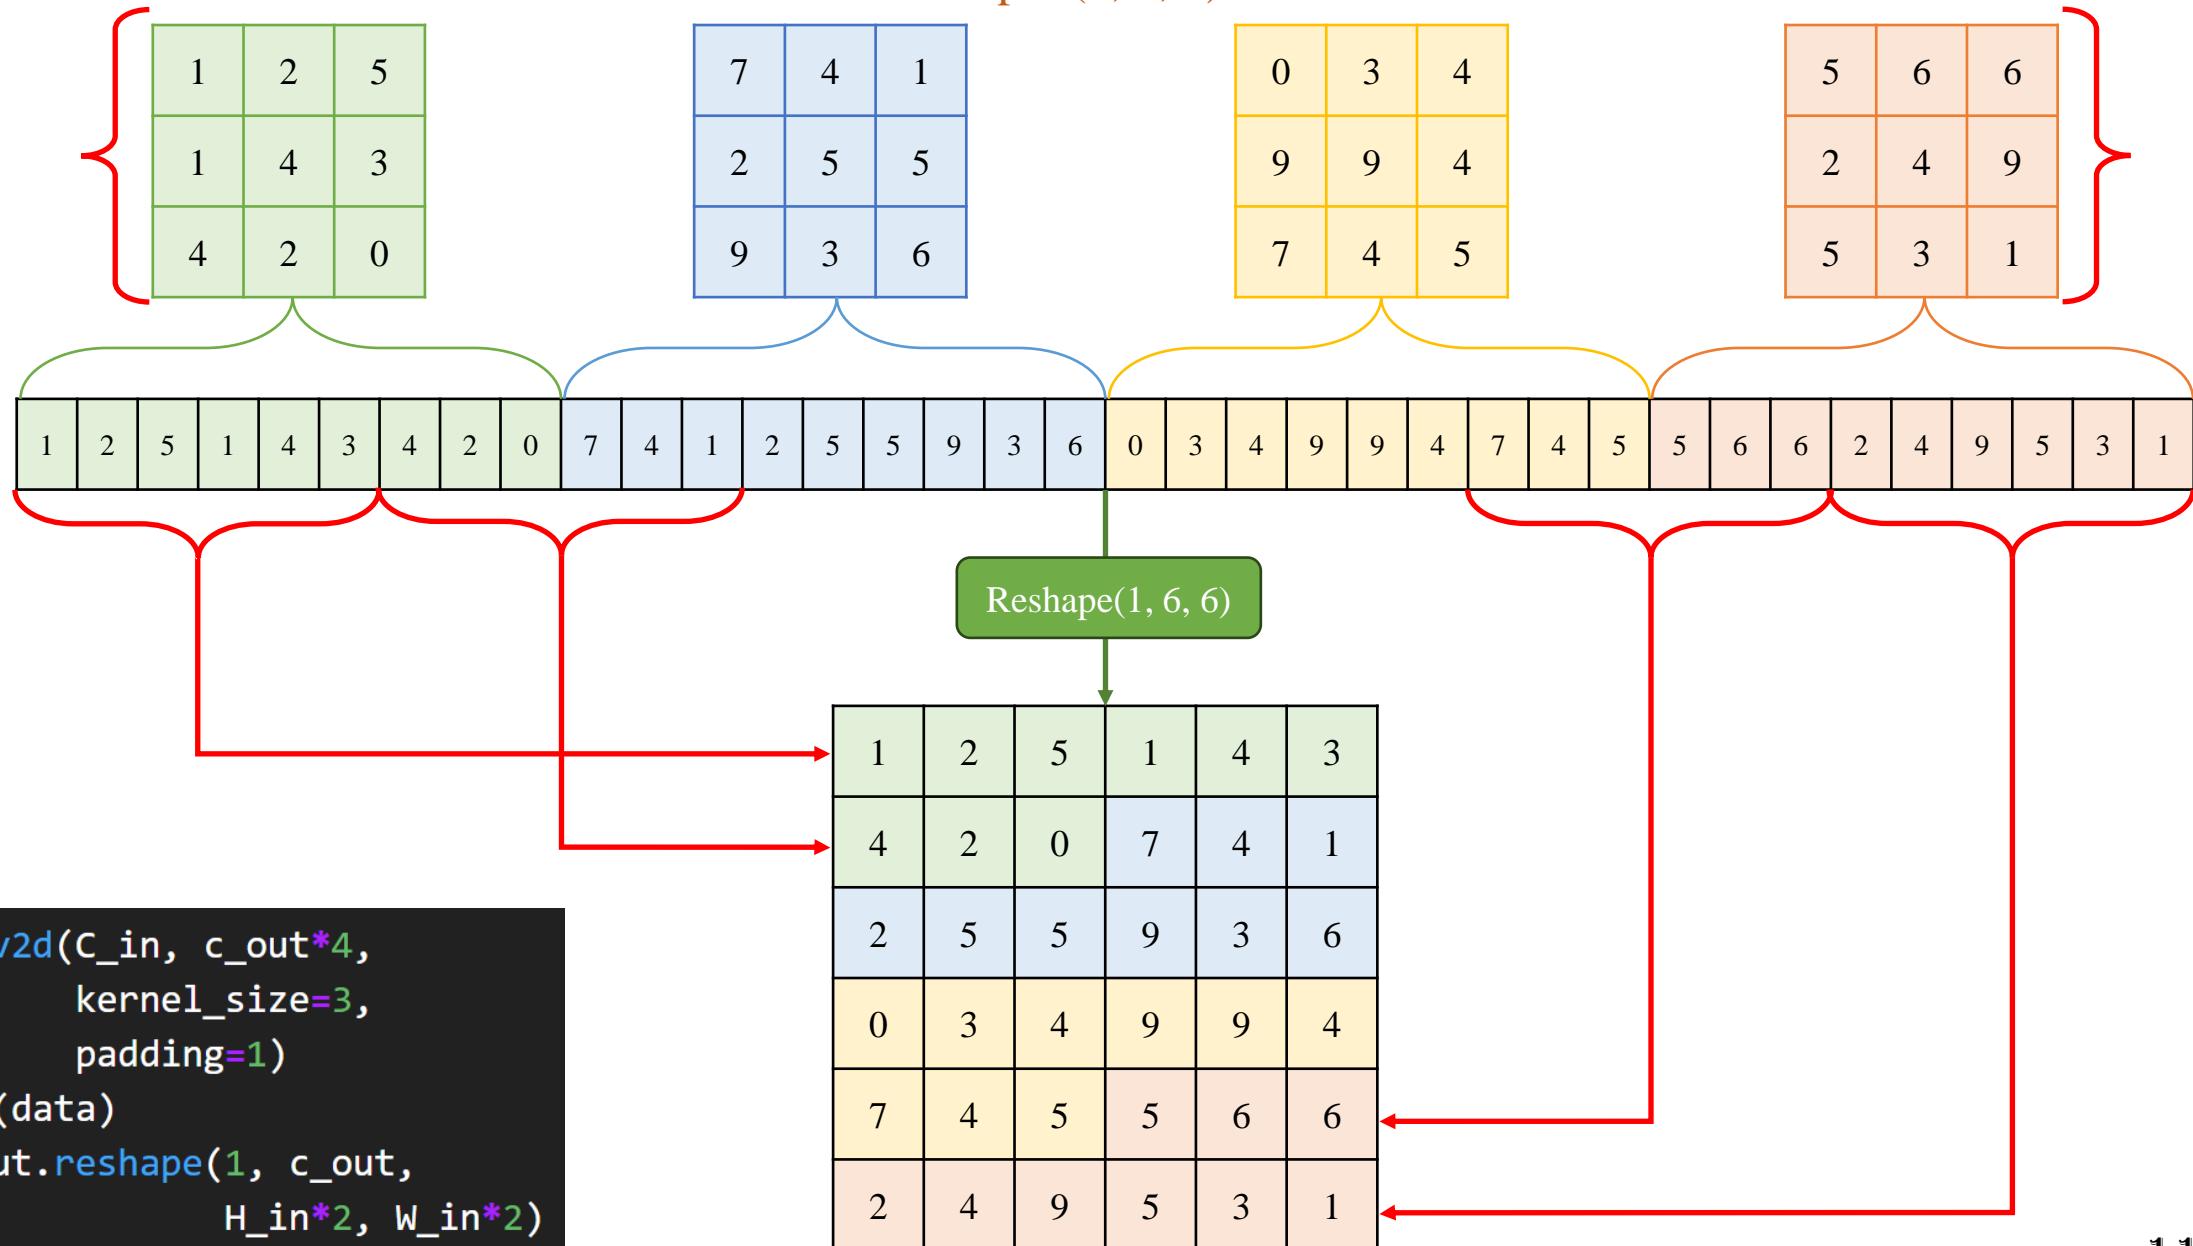
data =

1	5
4	8

1.0	2.0	4.0	5.0
1.7	2.7	4.7	5.7
3.2	4.2	6.2	7.2
4.0	5.0	7.0	8.0

Bilinear
interpolation

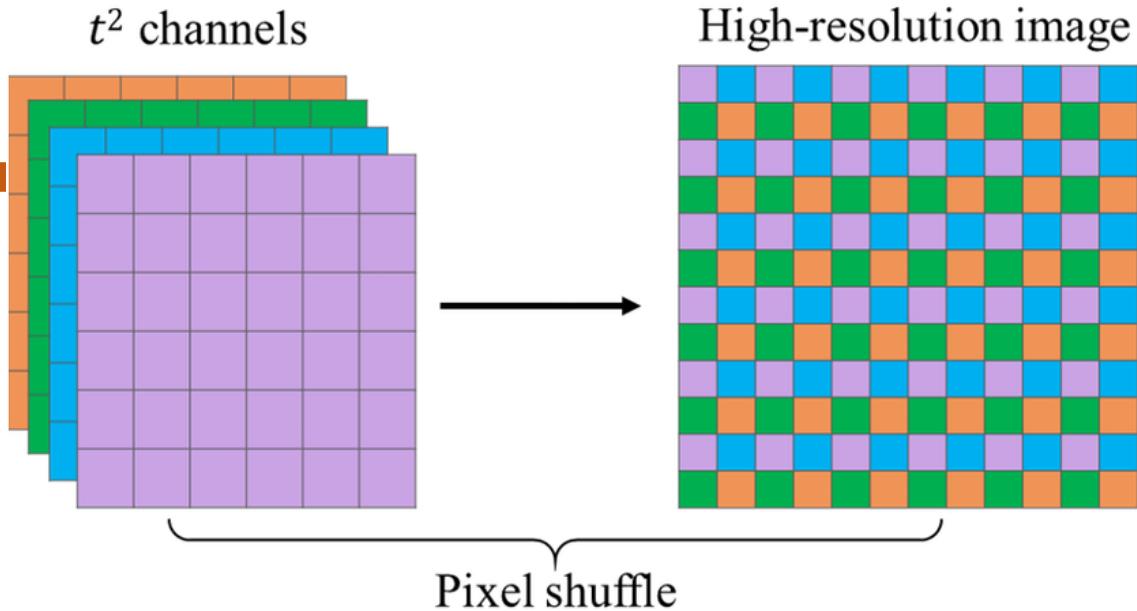
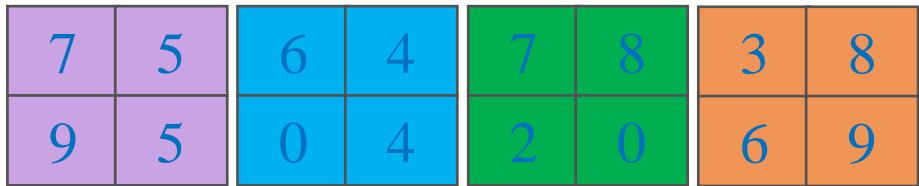
Reshape



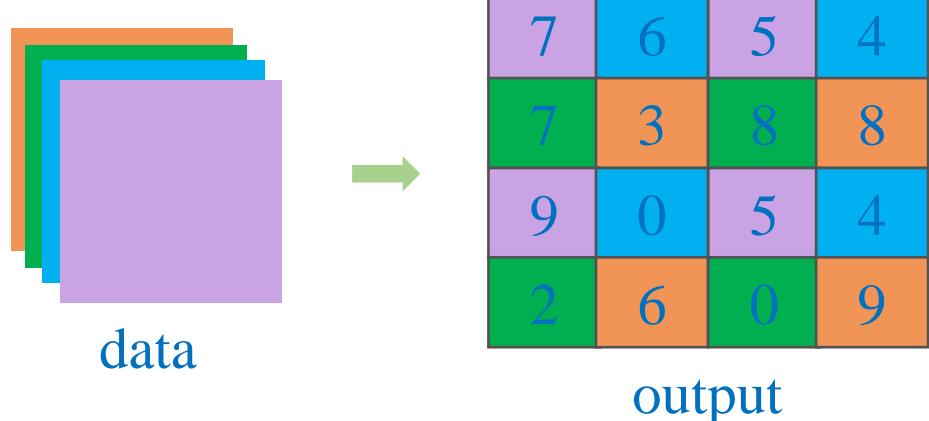
How to Increase Feature Map

❖ Pixel Shuffle

$$(B, C \times r^2, H, W) \rightarrow (B, C, H \times r, W \times r)$$



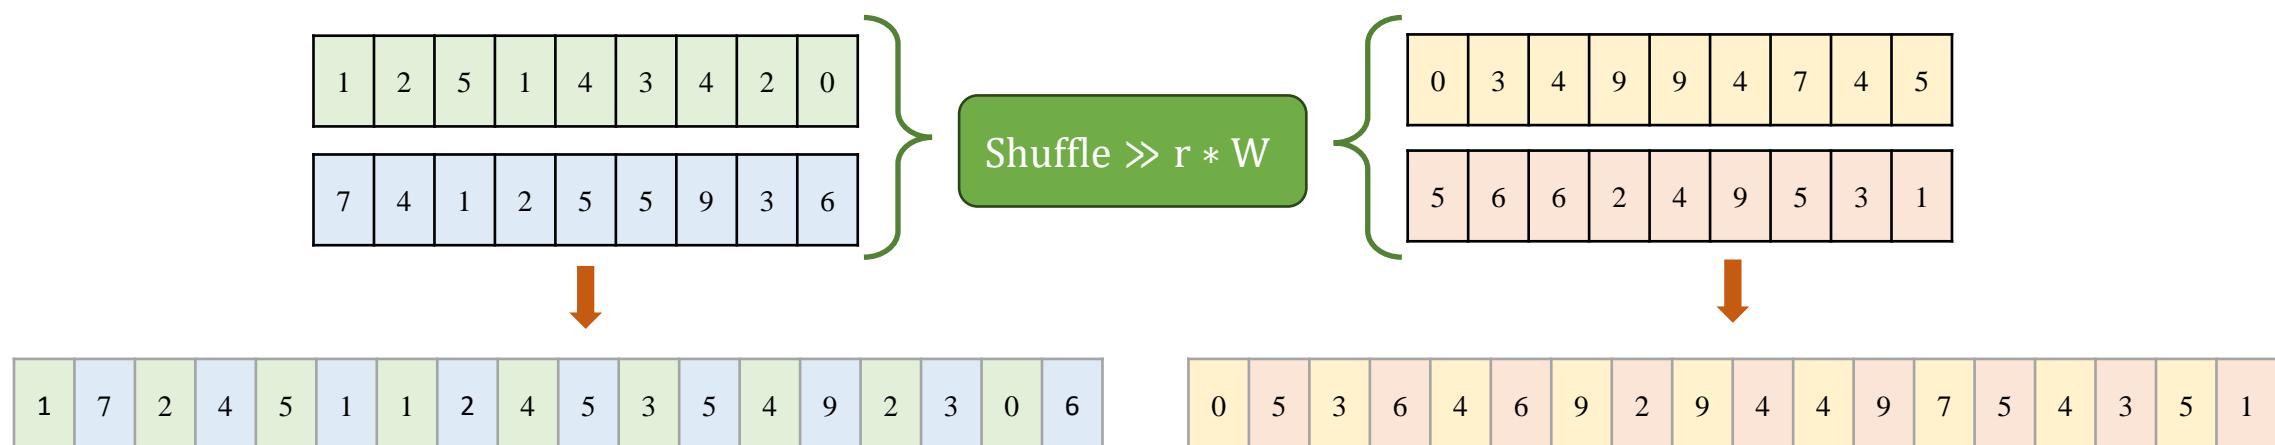
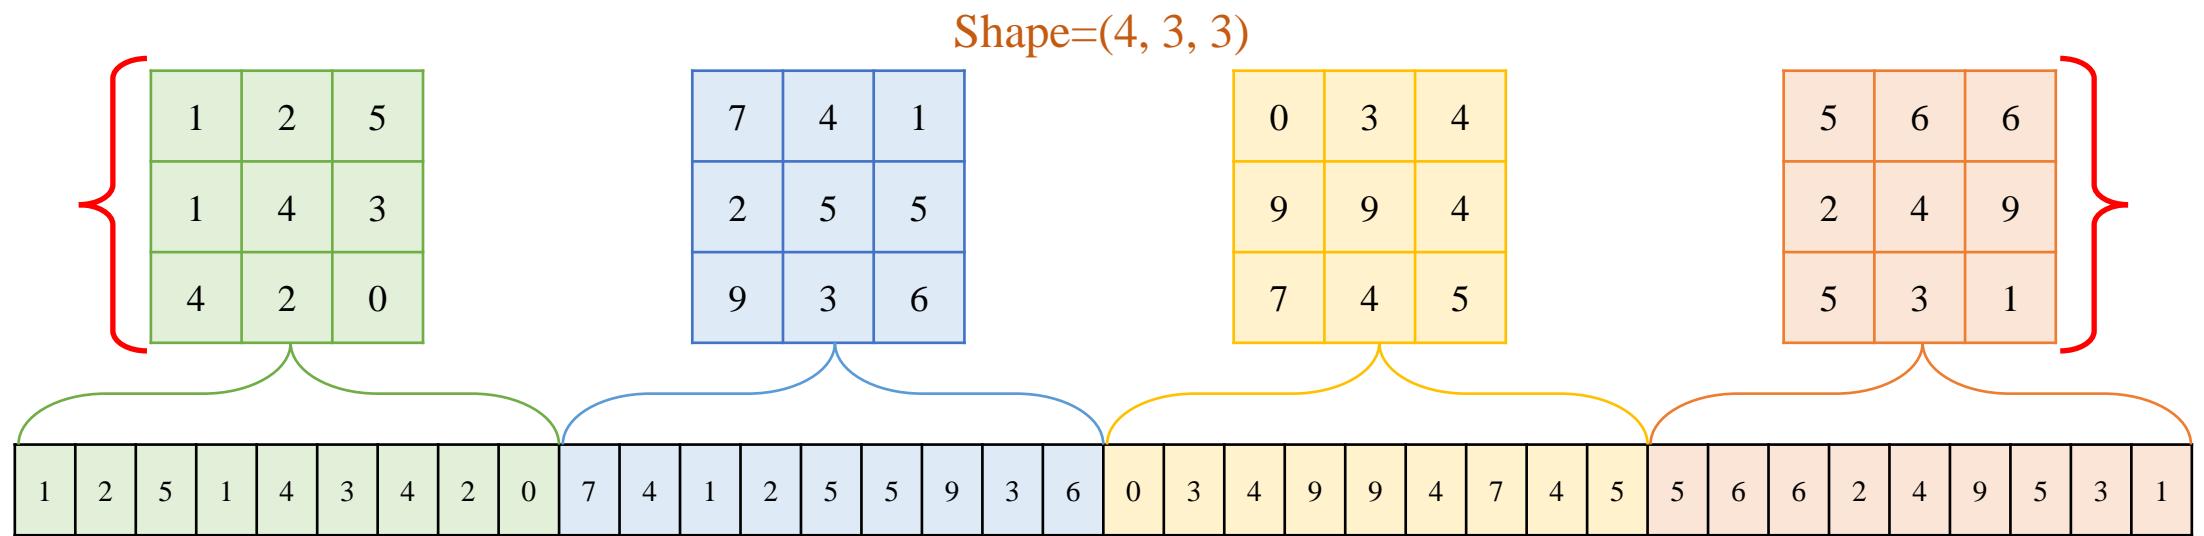
https://www.researchgate.net/figure/The-pixel-shuffle-layer-transforms-feature-maps-from-the-LR-domain-to-the-HR-image_fig3_339531308



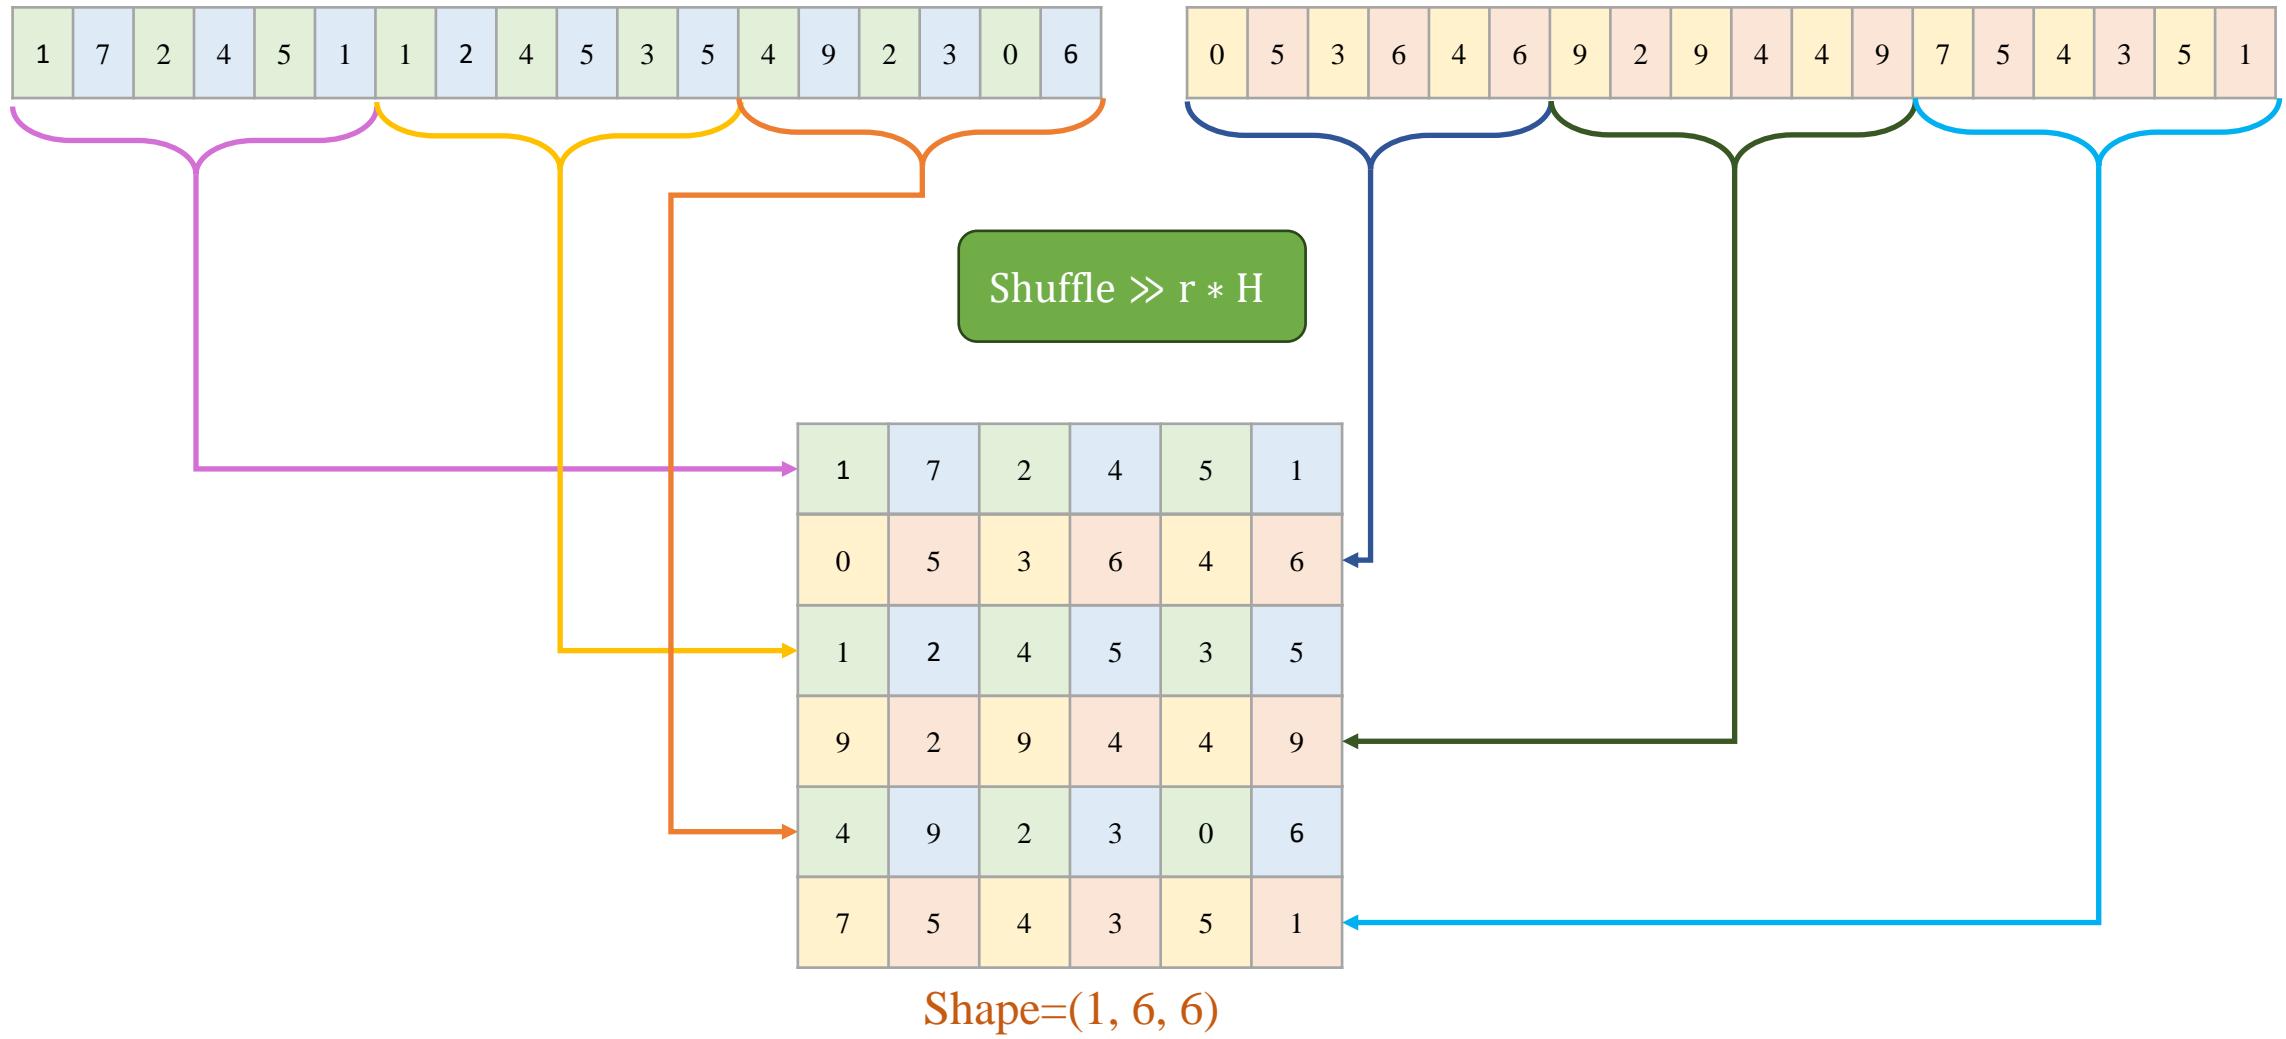
```
decoder = nn.Sequential(  
    nn.Conv2d(in_channels,  
             out_channels*4,  
             kernel_size=3,  
             padding=1),  
    # ...  
    nn.PixelShuffle(2))  
  
# forward  
x = decoder(x)
```

PixelShuffle

$$(C * r^2, W, H) \longrightarrow (C, r * W, r * H)$$



PixelShuffle



Unpooling

1	2	3	4
5	6	7	8
9	8	7	6
5	4	3	2

Data 1

Data 2

6	8
9	7



[[5, 7],
 [8, 10]]

0	0	0	0
0	6	0	8
9	0	7	0
0	0	0	0

Data 3

```
data = torch.Tensor([[1, 2, 3, 4],  
                    [5, 6, 7, 8],  
                    [9, 8, 7, 6],  
                    [5, 4, 3, 2]])  
  
data = data.reshape(1, 1, 4, 4)  
  
mpool = nn.MaxPool2d(2, return_indices=True)  
output_mp, indices = mpool(data)  
  
unpooling = nn.MaxUnpool2d(kernel_size=2)  
output_mup = unpooling(output_mp, indices)
```

```
# declaration  
encoder = nn.Sequential(  
    nn.Conv2d(in_channels, out_channels,  
             kernel_size=3, padding=1),  
    # ...  
    nn.MaxPool2d(2, return_indices=True))  
  
decoder = nn.Sequential(  
    nn.MaxUnpool2d(kernel_size=2),  
    # ...  
    nn.Conv2d(in_channels, out_channels,  
             kernel_size=3, padding=1))  
  
# forward  
x, indices = encoder(x)  
x = decoder(x, indices)
```


Convolution Transpose

❖ Example 1

tensorflow

```
Conv2DTranspose(filters=1, kernel_size=2, strides=1, padding='valid')
```

pytorch

```
ConvTranspose2d(in_channels=1, out_channels=1, kernel_size=2, stride=1, padding=0)
```

Stride S = 1

Padding p = 0

$$\text{data} = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$$

$D = 2$

$$\text{kernel} = \begin{bmatrix} 0.1 & -0.2 \\ 0.3 & -0.1 \end{bmatrix}$$

$K = 2$

Transpose padding
 $p' = K - 1 = 1$



$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$$



$$\begin{bmatrix} 0 & 0 & 0 & 0 \\ 0 & 1 & 2 & 0 \\ 0 & 4 & 3 & 0 \\ 0 & 0 & 0 & 0 \end{bmatrix}$$

$$\begin{bmatrix} 0 & 0 & 0 & 0 \\ 0 & 1 & 2 & 0 \\ 0 & 4 & 3 & 0 \\ 0 & 0 & 0 & 0 \end{bmatrix}$$

$$* \begin{bmatrix} 0.1 & -0.2 \\ 0.3 & -0.1 \end{bmatrix}$$

Output

$$= \begin{bmatrix} 0.1 & 0.0 & -0.4 \\ 0.6 & 0.3 & -1.0 \\ 0.9 & 0.9 & -0.4 \end{bmatrix}$$

Convolution Transpose

❖ Example 1

```
data_v1 = np.pad(data, 1)
print(data_v1)
```

```
[[0. 0. 0. 0.]
 [0. 1. 2. 0.]
 [0. 3. 4. 0.]
 [0. 0. 0. 0.]]
```

```
kernel = np.array([[ 0.1, -0.2],
                  [ 0.3, -0.1]])
print(kernel)
```

```
[[ 0.1 -0.2]
 [ 0.3 -0.1]]
```

```
from scipy import signal
signal.convolve2d(data_v1, kernel, mode='valid')
```

```
array([[ 0.1,  0. , -0.4],
       [ 0.6,  0.3, -1. ],
       [ 0.9,  0.9, -0.4]])
```

Numpy

0	0	0	0
0	1	2	0
0	4	3	0
0	0	0	0

$$\begin{matrix} & \begin{matrix} 0.1 & -0.2 \\ 0.3 & -0.1 \end{matrix} & = & \begin{matrix} 0.1 & 0.0 & -0.4 \\ 0.6 & 0.3 & -1.0 \\ 0.9 & 0.9 & -0.4 \end{matrix} \end{matrix}$$

$$O = D + (K - 1) = 3$$

```
conv_transpose = torch.nn.ConvTranspose2d(in_channels=1, out_channels=1,
                                         kernel_size=2, stride=1, padding=0)
```

```
output = conv_transpose(data)
print(output[0,0,:,:])
```

```
tensor([[ 0.1000,  0.0000, -0.4000],
        [ 0.6000,  0.3000, -1.0000],
        [ 0.9000,  0.9000, -0.4000]], grad_fn=<SliceBackward0>)
```

Pytorch

```
data = np.array([[1.0, 2.0],
                [3.0, 4.0]])
data = data.reshape(1, 2, 2, 1)
```

```
conv_transpose = tf.keras.layers.Conv2DTranspose(filters=1, kernel_size=2,
                                                 strides=1, padding='valid')
```

Tensorflow

```
output = conv_transpose(data)
print(output[0,:,:,:])
```

```
tf.Tensor(
[[ 0.1          0.           -0.4         ]
 [ 0.6          0.29999998 -1.          ]
 [ 0.90000004  0.90000004 -0.4         ]], shape=(3, 3), dtype=float32)
```

Convolution Transpose: Example 1 (Different view)

$$\text{data} = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$$

$$\text{kernel} = \begin{bmatrix} 0.1 & -0.2 \\ 0.3 & -0.1 \end{bmatrix}$$

Stride S = 1
Padding p = 0

$$1 \times \begin{bmatrix} 0.1 & -0.2 \\ 0.3 & -0.1 \end{bmatrix}$$

$$2 \times \begin{bmatrix} 0.1 & -0.2 \\ 0.3 & -0.1 \end{bmatrix}$$

$$3 \times \begin{bmatrix} 0.1 & -0.2 \\ 0.3 & -0.1 \end{bmatrix}$$

$$4 \times \begin{bmatrix} 0.1 & -0.2 \\ 0.3 & -0.1 \end{bmatrix}$$



$$\begin{bmatrix} 0.1 & -0.2 \\ 0.3 & -0.1 \end{bmatrix}$$

$$\begin{bmatrix} 0.2 & -0.4 \\ 0.6 & -0.2 \end{bmatrix}$$

$$\begin{bmatrix} 0.3 & -0.6 \\ 0.9 & -0.3 \end{bmatrix}$$

$$\begin{bmatrix} 0.4 & -0.8 \\ 1.2 & -0.4 \end{bmatrix}$$



$$\begin{bmatrix} 0.1 & -0.2 & \\ 0.3 & -0.1 & \\ & & \end{bmatrix}$$

$$+ \begin{bmatrix} & 0.2 & -0.4 \\ & 0.6 & -0.2 \\ & & \end{bmatrix}$$

$$+ \begin{bmatrix} & & \\ 0.3 & -0.6 & \\ 0.9 & -0.3 & \end{bmatrix}$$

$$+ \begin{bmatrix} & & \\ & 0.4 & -0.8 \\ & 1.2 & -0.4 \end{bmatrix}$$

$$= \begin{bmatrix} 0.1 & 0.0 & -0.4 \\ 0.6 & 0.3 & -1.0 \\ 0.9 & 0.9 & -0.4 \end{bmatrix}$$

Convolution Transpose

❖ Example 2

Padding $p = 0$
(padding='valid')

$$(D - K) \% S == 0$$

data =

1	2
3	4

$D = 2$

kernel =

0.1	-0.2
0.3	-0.1

$K = 2$

With stride $S = 2$

1	2
3	4



1	0	2
0	0	0
4	0	3



Transpose padding
 $p' = K - 1 = 1$

0	0	0	0	0
0	1	0	2	0
0	0	0	0	0
0	4	0	3	0
0	0	0	0	0

*

0.1	-0.2
0.3	-0.1

=

0.1	-0.2	0.2	-0.4
0.3	-0.1	0.6	-0.2
0.3	-0.6	0.4	-0.8
0.9	-0.3	1.2	-0.4

Convolution Transpose

❖ Example 2

```
data_v3 = np.insert(data, obj=[1], values=0, axis=1)
data_v3 = np.insert(data_v3, obj=[1], values=0, axis=0)
print(data_v3)
```

```
[[1. 0. 2.]
 [0. 0. 0.]
 [3. 0. 4.]]
```

```
data_v3 = np.pad(data_v3, 1)
print(data_v3)
```

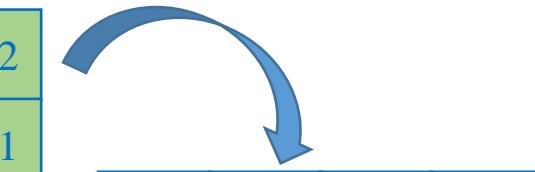
```
[[0. 0. 0. 0. 0.]
 [0. 1. 0. 2. 0.]
 [0. 0. 0. 0. 0.]
 [0. 3. 0. 4. 0.]
 [0. 0. 0. 0. 0.]]
```

```
signal.convolve2d(data_v3, kernel, mode='valid')
```

```
array([[ 0.1, -0.2,  0.2, -0.4],
       [ 0.3, -0.1,  0.6, -0.2],
       [ 0.3, -0.6,  0.4, -0.8],
       [ 0.9, -0.3,  1.2, -0.4]])
```

Numpy

0	0	0	0	0.1	-0.2
0	1	0	2	0.3	-0.1
0	0	0	0	0	
0	4	0	3	0	
0	0	0	0	0	



0.1	-0.2	0.2	-0.4
0.3	-0.1	0.6	-0.2
0.3	-0.6	0.4	-0.8
0.9	-0.3	1.2	-0.4

```
#torch.nn.ConvTranspose2d(in_channels=C, out_channels=1,
#                           kernel_size=(2,2), stride=2, padding=0)

output = conv_transpose(data)
```

Pytorch

```
tensor([[ 0.1000, -0.2000,  0.2000, -0.4000],
        [ 0.3000, -0.1000,  0.6000, -0.2000],
        [ 0.3000, -0.6000,  0.4000, -0.8000],
        [ 0.9000, -0.3000,  1.2000, -0.4000]], grad_fn=<SliceBackward0>)
```

```
conv_transpose = Conv2DTranspose(filters=1, kernel_size=2,
                                  strides=2, padding='valid')

output = conv_transpose(data)

tf.Tensor(
[[ 0.1           -0.2            0.2            -0.4          ],
 [ 0.3           -0.1            0.6            -0.2          ],
 [ 0.3           -0.6            0.4            -0.8          ],
 [ 0.90000004   -0.3            1.2            -0.4          ]], shape=(4, 4),
```

Tensorflow

Convolution Transpose

Example 2 (Different view)

$$\text{data} = \begin{array}{|c|c|} \hline 1 & 2 \\ \hline 3 & 4 \\ \hline \end{array}$$

$$\text{kernel} = \begin{array}{|c|c|} \hline 0.1 & -0.2 \\ \hline 0.3 & -0.1 \\ \hline \end{array}$$

Padding $p = 0$
 $(D - K) \% S == 0$

With stride $S = 2$

$$1 \times \begin{array}{|c|c|} \hline 0.1 & -0.2 \\ \hline 0.3 & -0.1 \\ \hline \end{array}$$

$$2 \times \begin{array}{|c|c|} \hline 0.1 & -0.2 \\ \hline 0.3 & -0.1 \\ \hline \end{array}$$

$$3 \times \begin{array}{|c|c|} \hline 0.1 & -0.2 \\ \hline 0.3 & -0.1 \\ \hline \end{array}$$

$$4 \times \begin{array}{|c|c|} \hline 0.1 & -0.2 \\ \hline 0.3 & -0.1 \\ \hline \end{array}$$

$$\begin{array}{|c|c|} \hline 0.1 & -0.2 \\ \hline 0.3 & -0.1 \\ \hline \end{array}$$

$$\begin{array}{|c|c|} \hline 0.2 & -0.4 \\ \hline 0.6 & -0.2 \\ \hline \end{array}$$

$$\begin{array}{|c|c|} \hline 0.3 & -0.6 \\ \hline 0.9 & -0.3 \\ \hline \end{array}$$

$$\begin{array}{|c|c|} \hline 0.4 & -0.8 \\ \hline 1.2 & -0.4 \\ \hline \end{array}$$

$$\begin{array}{|c|c|c|c|} \hline 0.1 & -0.2 & & \\ \hline 0.3 & -0.1 & & \\ \hline & & & \\ \hline & & & \\ \hline \end{array}$$

$$+ \begin{array}{|c|c|c|c|} \hline & & 0.2 & -0.4 \\ \hline & & 0.6 & -0.2 \\ \hline & & & \\ \hline & & & \\ \hline \end{array}$$

$$+ \begin{array}{|c|c|c|c|} \hline & & & \\ \hline & & & \\ \hline 0.3 & -0.6 & & \\ \hline 0.9 & -0.3 & & \\ \hline & & & \\ \hline & & & \\ \hline \end{array}$$

$$+ \begin{array}{|c|c|c|c|} \hline & & & \\ \hline & & & \\ \hline & & 0.4 & -0.8 \\ \hline & & 1.2 & -0.4 \\ \hline & & & \\ \hline & & & \\ \hline \end{array}$$

$$= \begin{array}{|c|c|c|c|} \hline 0.1 & -0.2 & 0.2 & -0.4 \\ \hline 0.3 & -0.1 & 0.6 & -0.2 \\ \hline 0.3 & -0.6 & 0.4 & -0.8 \\ \hline 0.9 & -0.3 & 1.2 & -0.4 \\ \hline \end{array}$$

ConvTranspose

Input

1.	2.
3.	4.

Shape=(2,2)

1.

3.	0.	-2.
1.	1.	1.
-3.	0.	2.

=

3.	0.	-2.
1.	1.	1.
-3.	0.	2.

2.

3.	0.	-2.
1.	1.	1.
-3.	0.	2.

=

6.	0.	-4.
2.	2.	2.
-6.	0.	4.

Kernel

3.	0.	-2.
1.	1.	1.
-3.	0.	2.

Shape=(3,3)

3.

3.	0.	-2.
1.	1.	1.
-3.	0.	2.

=

9.	0.	-6.
3.	3.	3.
-9.	0.	6.

4.

3.	0.	-2.
1.	1.	1.
-3.	0.	2.

=

12.	0.	-8.
4.	4.	4.
-12.	0.	8.

3.	0.	-2.	6.	0.	-4.
1.	1.	1.	2.	2.	2.
-3.	0.	2.	-6.	0.	4.
9.	0.	-6.	12.	0.	-8.

Stride = 2

9.	0.	-6.	12.	0.	-8.
3.	3.	3.	4.	4.	4.
-9.	0.	6.	-12.	0.	8.

Sum

3.	0.	4.	0.	-4.
1.	1.	3.	2.	2.
6.	0.	2.	0.	-4.
3.	3.	7.	4.	4.
-9.	0.	-6.	0.	8.

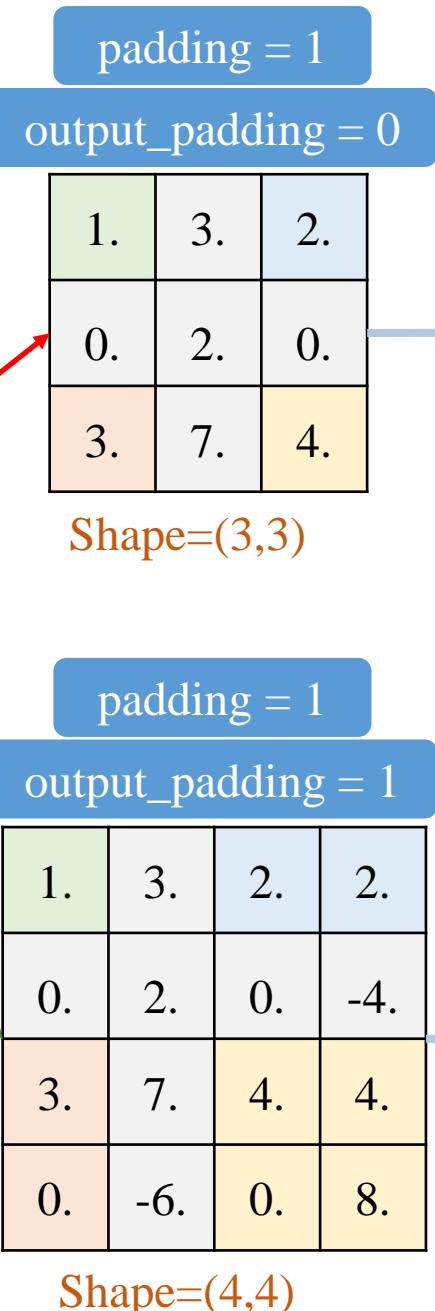
Shape=(5,5)

ConvTranspose

padding = 0
output_padding = 0

3.	0.	4.	0.	-4.
1.	1.	3.	2.	2.
6.	0.	2.	0.	-4.
3.	3.	7.	4.	4.
-9.	0.	-6.	0.	8.

Shape=(5,5)



$$O = S(D - 1) + (K - 1) - 2 \times P + P_o + 1$$

Bias
+0.1

Output

1.1	3.1	2.1
0.1	2.1	0.1
3.1	7.1	4.1

Shape=(3,3)

Bias
+0.1

Output

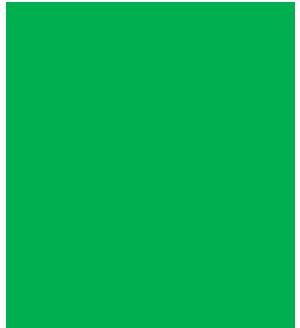
1.1	3.1	2.1	2.1
0.1	2.1	0.1	-3.9
3.1	7.1	4.1	4.1
0.1	-5.9	0.1	8.1

Shape=(4,4)

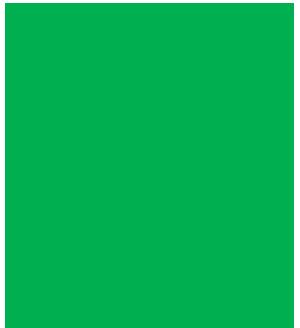
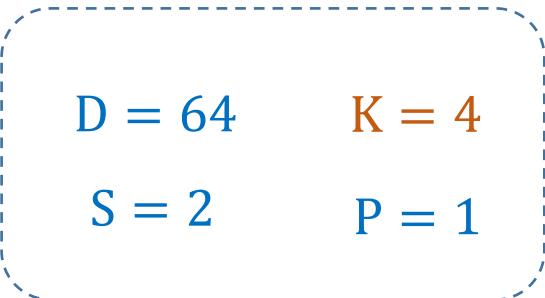
Convolution Transpose

```
nn.ConvTranspose2d(in_channels, out_channels,  
                 kernel_size=4, stride=2,  
                 padding=1, output_padding=0)
```

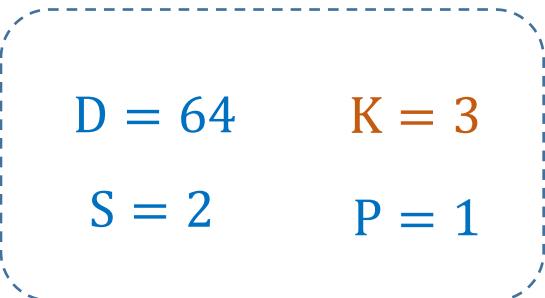
- ❖ To upsample feature maps 2x



64x64x1



64x64x1



Transpose padding: $P_o = 0$

Output size
 $O = S(D - 1) + (K - 1) - 2 \times P + P_o + 1$
 $= 2 * 63 + 3 - 2 + 1 + 1 = 128$

```
nn.ConvTranspose2d(in_channels, out_channels,  
                 kernel_size=3, stride=2,  
                 padding=1, output_padding=0)
```

Output size

$$O = S(D - 1) + (K - 1) - 2 \times P + P_o + 1
= 2 * 63 + 2 - 2 + 1 + 1 = 127$$

Further reading:

<https://arxiv.org/pdf/1603.07285v1.pdf>

Convolution Transpose

$$\begin{array}{ll} D = 64 & K = 3 \\ S = 2 & P = 1 \end{array}$$

$$\begin{array}{ll} D = 64 & K = 4 \\ S = 2 & P = 1 \end{array}$$

```
import torch
import numpy as np

N, H, W, C = 1, 64, 64, 1
data = np.random.randn(N, C, H, W)
data = torch.tensor(data, dtype=torch.float)

conv_transpose = torch.nn.ConvTranspose2d(in_channels=1, out_channels=1,
                                         kernel_size=(3,3), stride=2, padding=1)

output = conv_transpose(data)
print(output.shape)

torch.Size([1, 1, 127, 127])
```

```
import torch
import numpy as np

N, H, W, C = 1, 64, 64, 1
data = np.random.randn(N, C, H, W)
data = torch.tensor(data, dtype=torch.float)

conv_transpose = torch.nn.ConvTranspose2d(in_channels=1, out_channels=1,
                                         kernel_size=(4,4), stride=2, padding=1)

output = conv_transpose(data)
print(output.shape)

torch.Size([1, 1, 128, 128])
```

Further reading:

<https://arxiv.org/pdf/1603.07285v1.pdf>

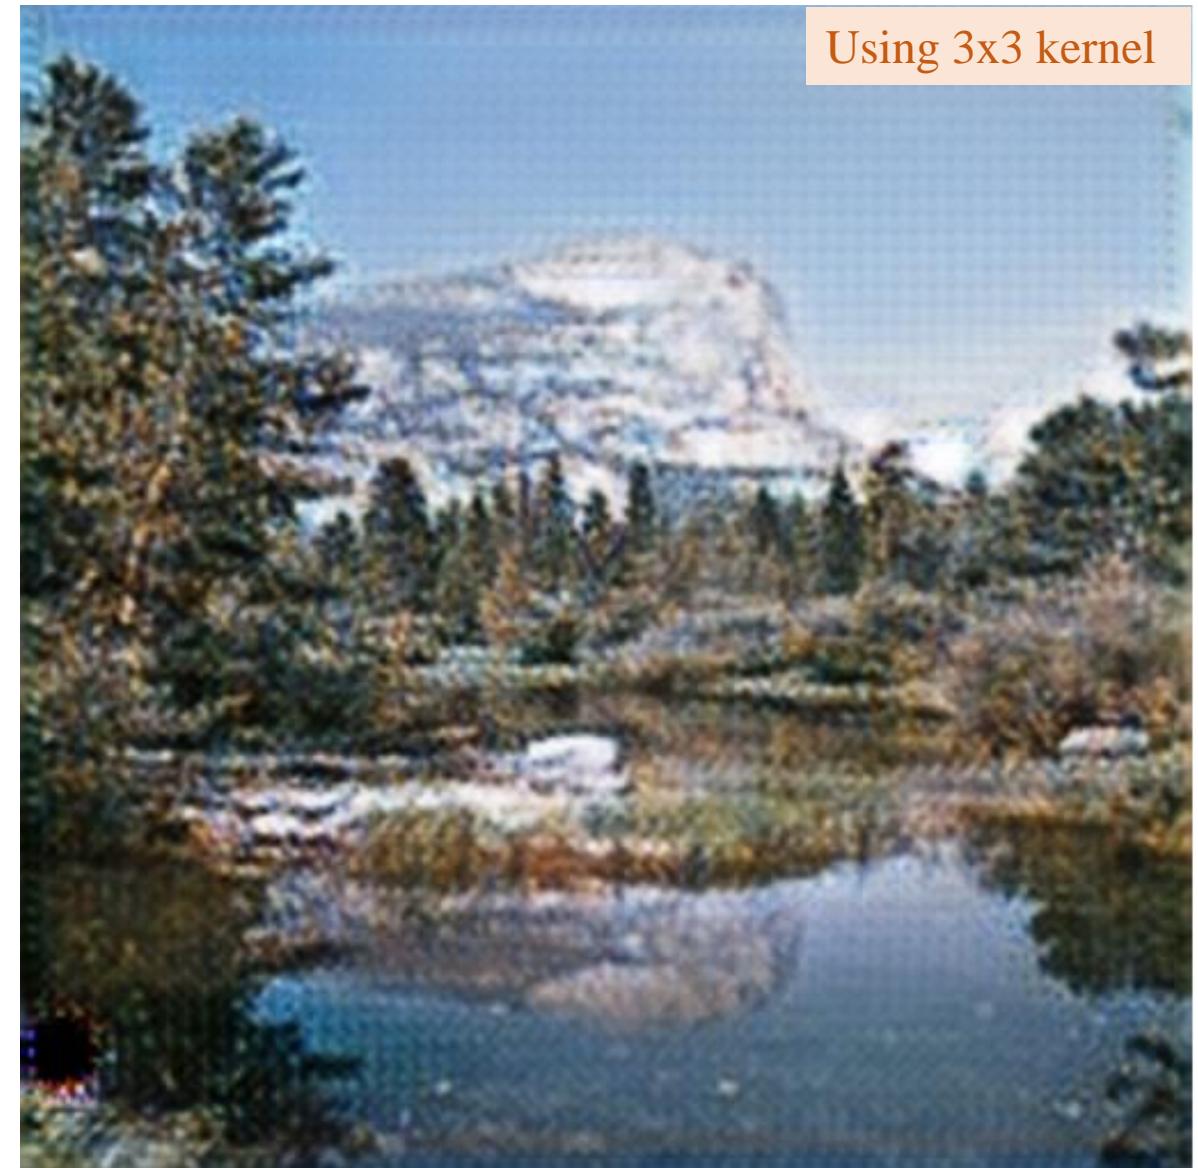
Using 3x3 or 4x4

❖ Convolution Transpose

```
nn.ConvTranspose2d(in_channels, out_channels,  
                  kernel_size=4, stride=2,  
                  padding=1, output_padding=0)
```

<https://github.com/junyanz/pytorch-CycleGAN-and-pix2pix/issues/78>

Using 3x3 kernel



Outline

- Segmentation (Unet Using Interpolation)
- Alternatives to Increase Feature Resolution
- Colorization
- Super-resolution
- Denoising Survey
- Super-resolution Survey

Image Denoising and Colorization

Noisy images



Model

Clean images



Grayscale images



Model

Color images



Noisy and grayscale images



Model

Clean and color images



Image Translation

❖ Super-resolution



Image Translation

❖ Edge2Scene

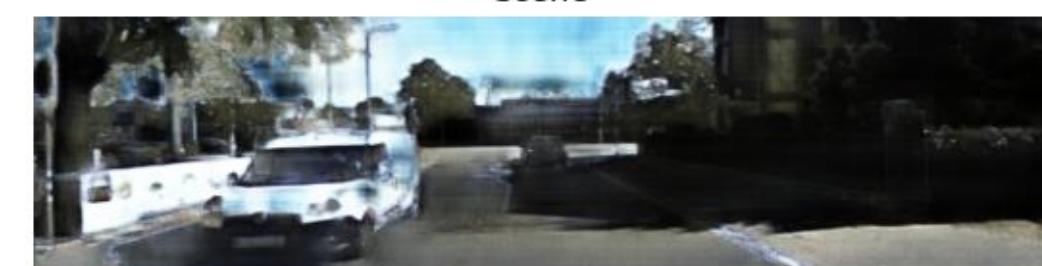
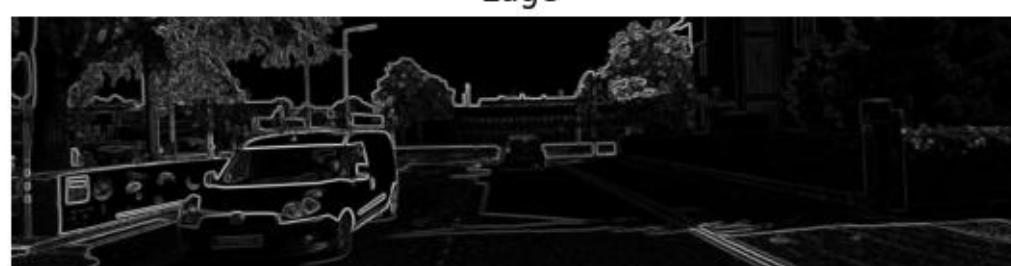
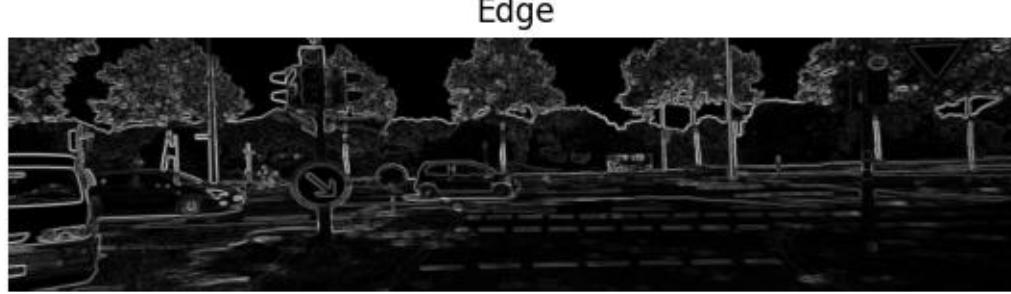


Image Translation

Deblur

1	Data preparation	3	Loss and optimizer
2	Network construction	4	Training



Input Images



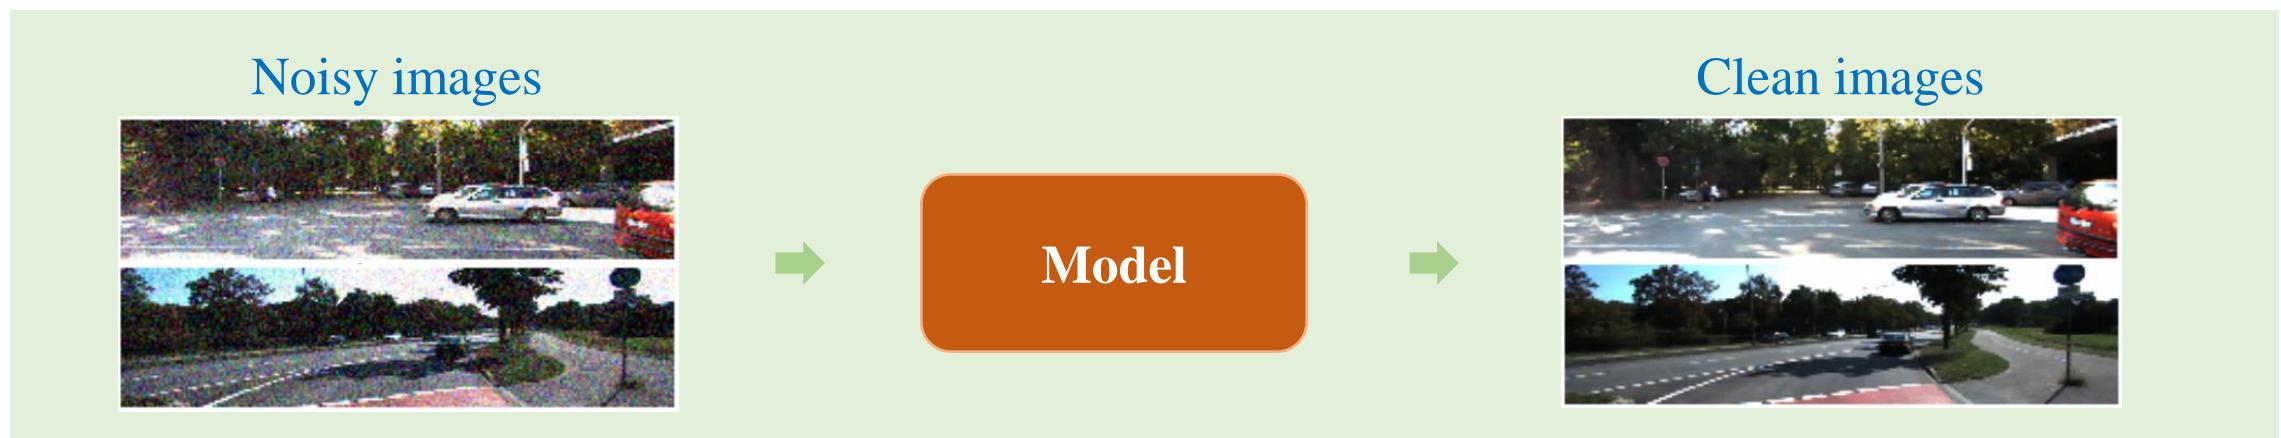
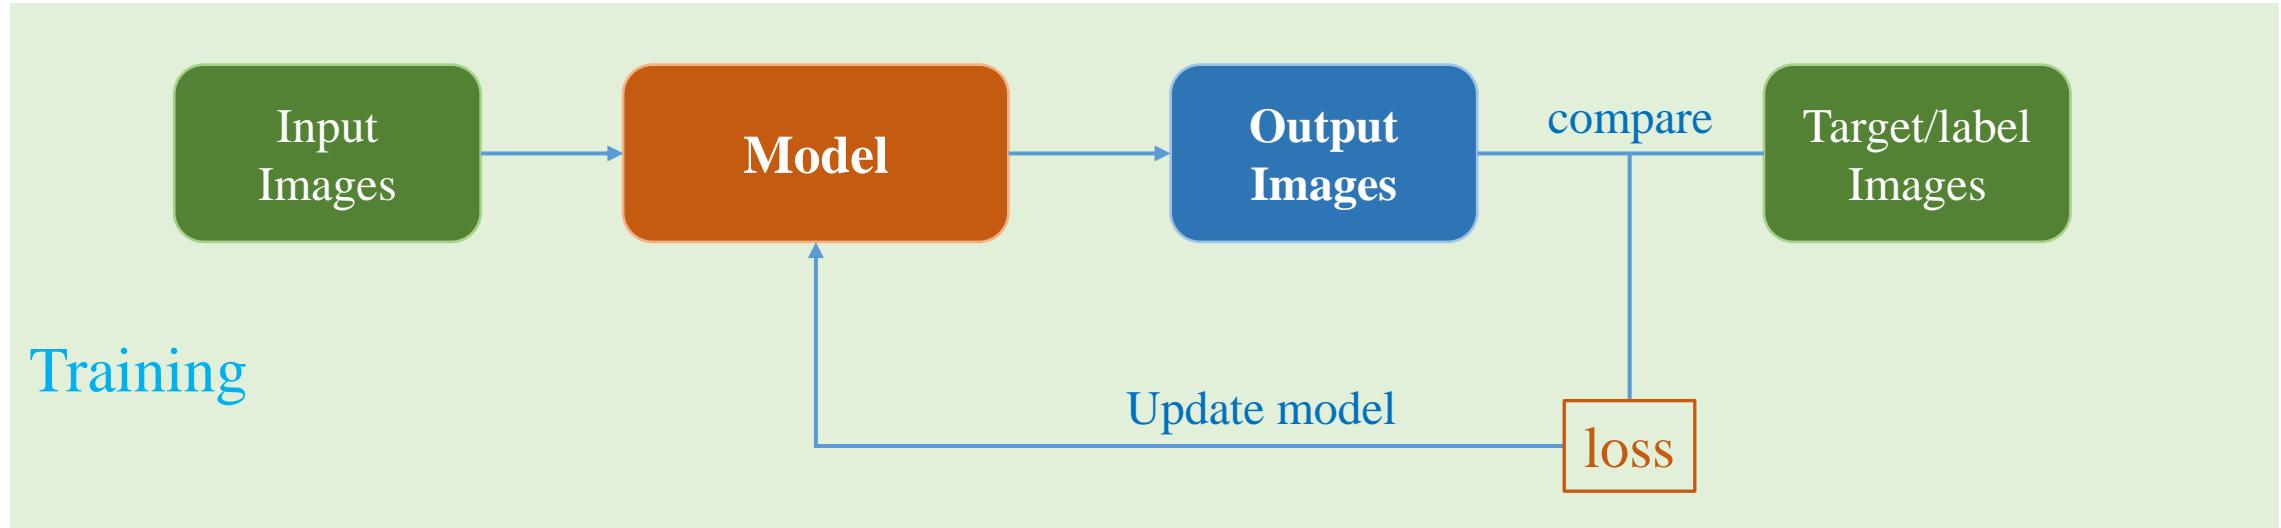
Target Images



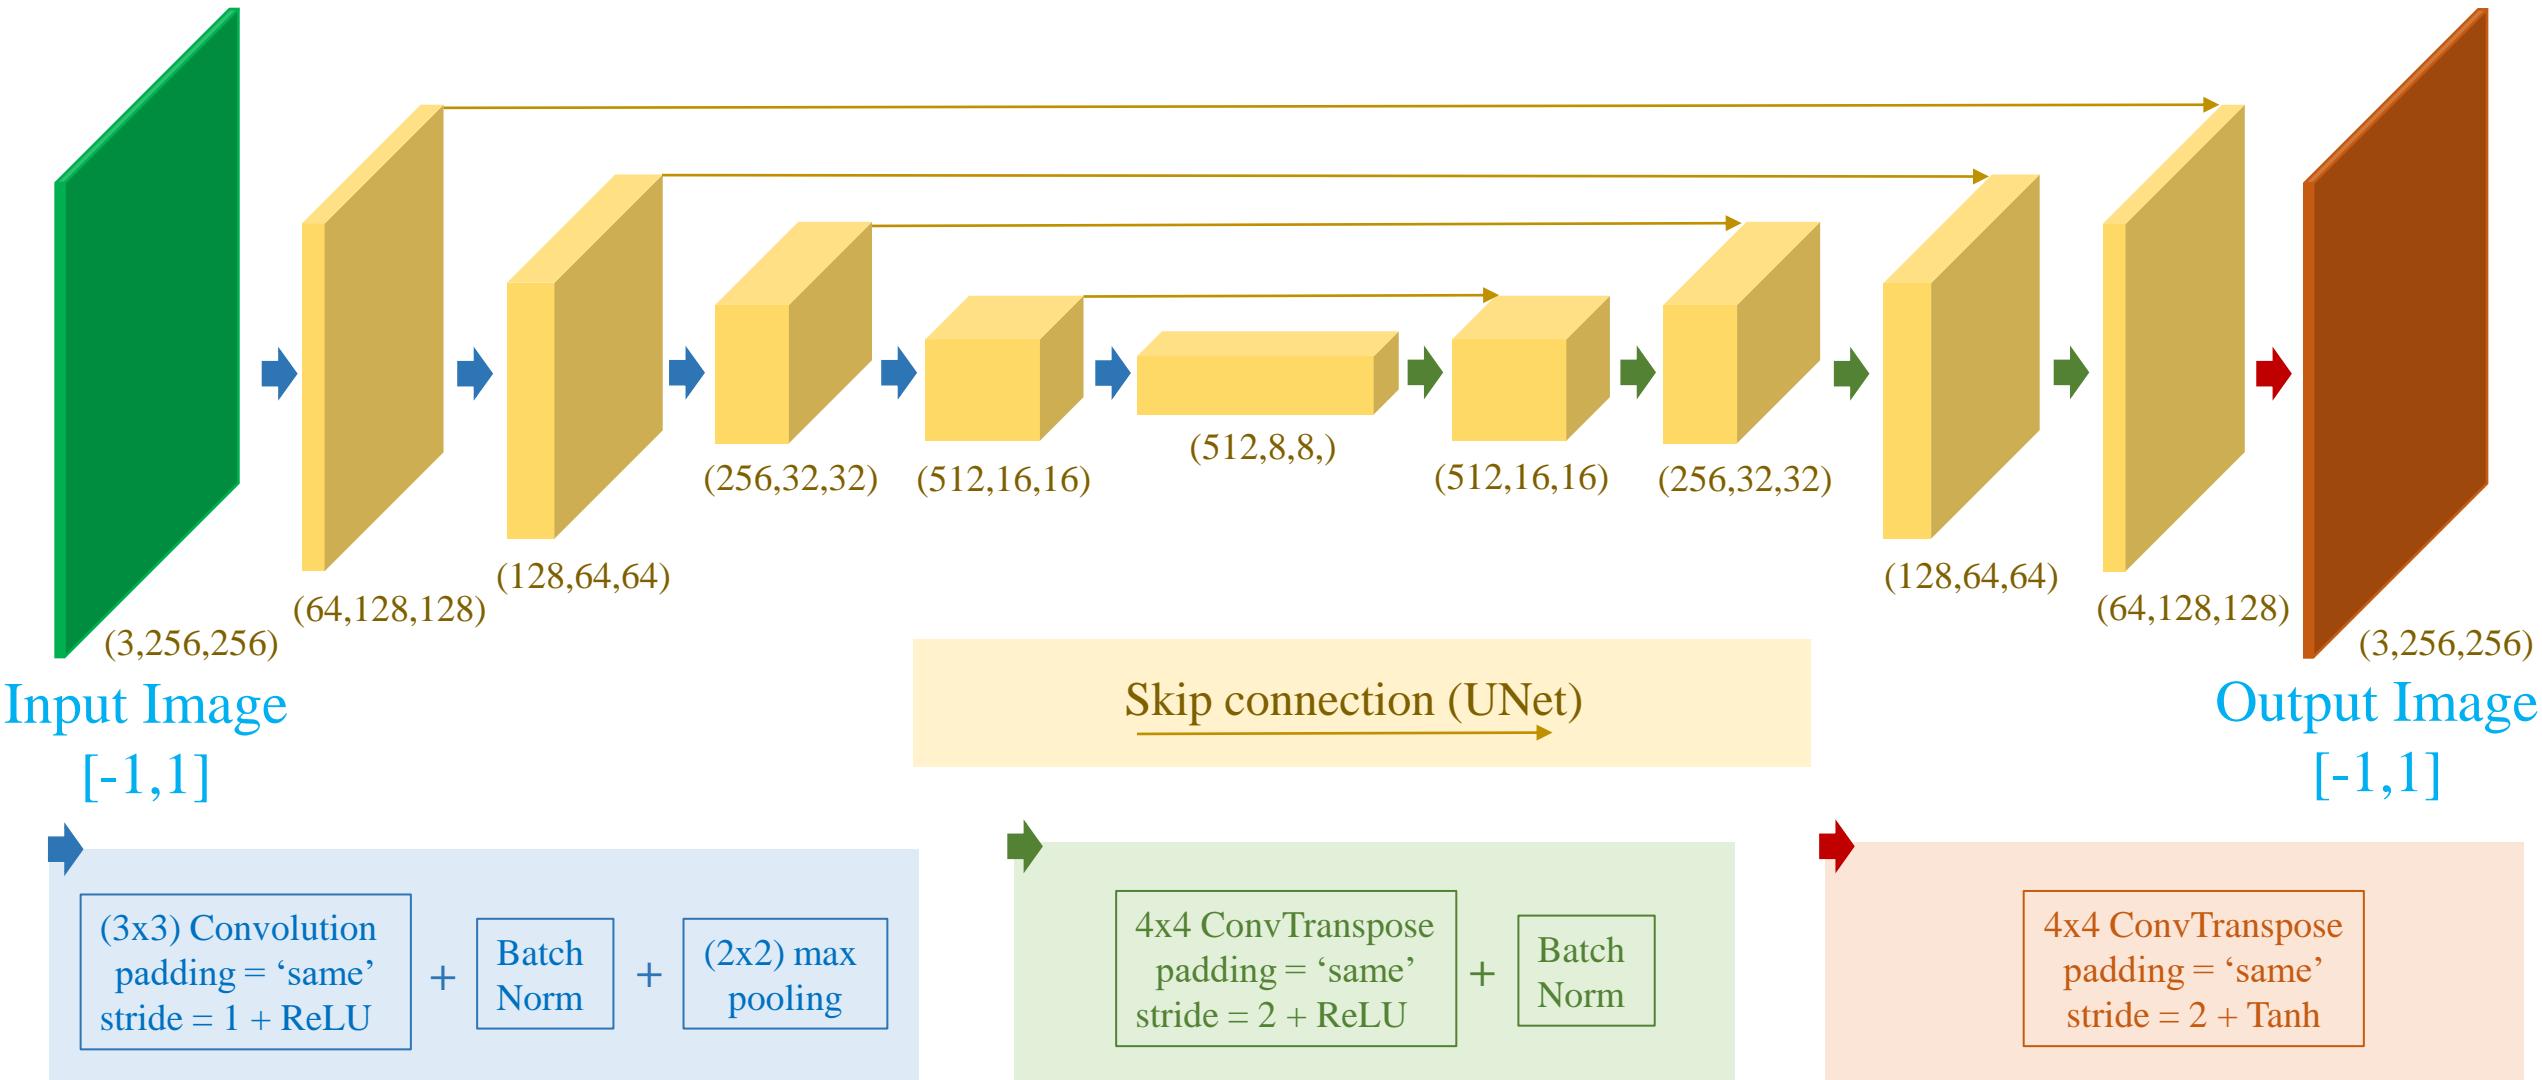
Predicted Images

Image Denoising and Colorization

❖ Flowchart



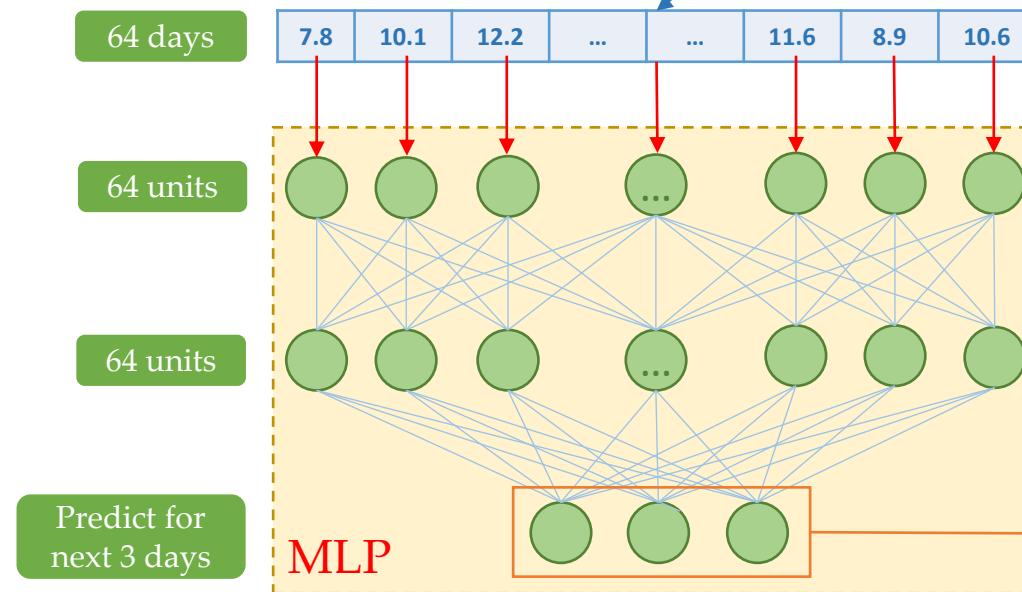
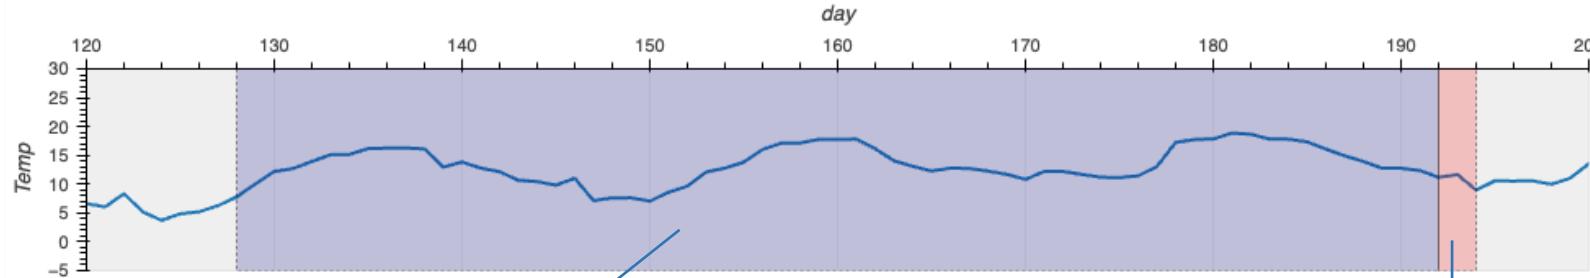
Denoising/Colorization/Deblur/...



Regression Problem

$$L1(I1, I2) = \frac{1}{H * W * C} \sum_{h=1}^H \sum_{w=1}^W \sum_{c=1}^C |I1_{hwc} - I2_{hwc}|$$

Input/Output + Loss Function

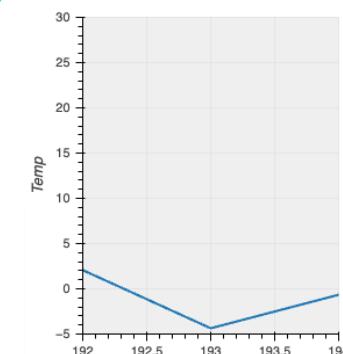
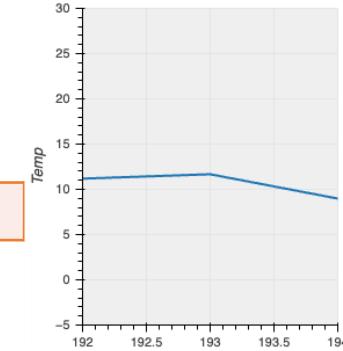


update

predict

MAE Loss

11.4



Further Reading

Unet-related Papers

U-Net: Convolutional Networks for Biomedical Image Segmentation

Olaf Ronneberger, Philipp Fischer, and Thomas Brox

Computer Science Department and BIOSS Centre for Biological Signalling Studies,
University of Freiburg, Germany
ronneber@informatik.uni-freiburg.de,
WWW home page: <http://lmb.informatik.uni-freiburg.de/>

TransUNet: Transformers Make Strong Encoders for Medical Image Segmentation

Jieneng Chen¹, Yongyi Lu¹, Qihang Yu¹, Xiangde Luo²,
Ehsan Adeli³, Yan Wang⁴, Le Lu⁵, Alan L. Yuille¹, and Yuyin Zhou³

¹Johns Hopkins University

²University of Electronic Science and Technology of China

³Stanford University

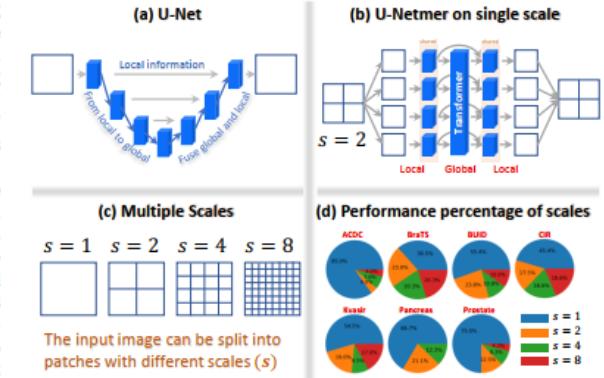
⁴East China Normal University

⁵PAII Inc.

U-Netmer: U-Net meets Transformer for medical image segmentation

Sheng He, Rina Bao, P. Ellen Grant, Yangming Ou

Abstract—The combination of the U-Net based deep learning models and Transformer is a new trend for medical image segmentation. U-Net can extract the detailed local semantic and texture information and Transformer can learn the long-rang dependencies among pixels in the input image. However, directly adapting the Transformer for segmentation has “token-flatten” problem (flattens the local patches into 1D tokens which loses the interaction among pixels within local patches) and “scale-sensitivity” problem (uses a fixed scale to split the input image into local patches). Compared to directly combining U-Net and Transformer, we propose a new global-local fashion combination of U-Net and Transformer, named U-Netmer, to solve the two problems. The proposed U-Netmer splits an input image into local patches. The global-context information among local patches is learnt by the self-attention mechanism in Transformer and U-Net segments each local patch instead of flattening into tokens to solve the ‘token-flatten’ problem. The U-Netmer can segment the input image with different patch sizes with the identical



3D TransUNet: Advancing Medical Image Segmentation through Vision Transformers

Jieneng Chen, Jieru Mei, Xianhang Li, Yongyi Lu, Qihang Yu, Qingyue Wei, Xiangde Luo, Yutong Xie, Ehsan Adeli, Yan Wang, Matthew Lungren, Lei Xing, Le Lu, Alan Yuille, Yuyin Zhou

Abstract—Medical image segmentation plays a crucial role in advancing healthcare systems for disease diagnosis and treatment planning. The u-shaped architecture, popularly known as U-Net, has proven highly successful for various medical image segmentation tasks. However, U-Net’s convolution-based operations inherently limit its ability to model long-range dependencies effectively.

To address these limitations, researchers have turned to Transformers, renowned for their global self-attention mechanisms, as alternative architectures. One popular network is our previous TransUNet, which leverages Transformers’ self-attention to complement U-Net’s localized information with the global context. In this paper, we extend

encoder and decoder into the u-shaped medical image segmentation architecture. TransUNet outperforms competitors in various medical applications, including multi-organ segmentation, pancreatic tumor segmentation, and hepatic vessel segmentation. It notably surpasses the top solution in the BrasTS2021 challenge. Code and models are available at <https://github.com/Beckschen/3D-TransUNet>.

I. INTRODUCTION

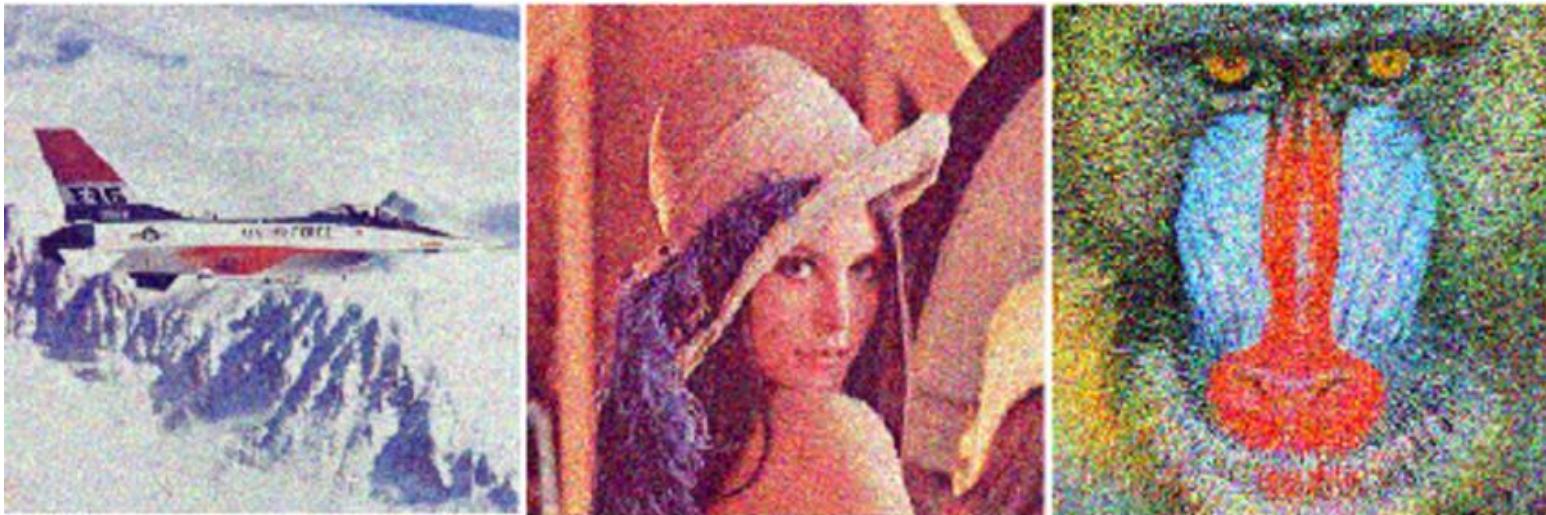
Outline

- Segmentation (Unet Using Interpolation)
- Alternatives to Increase Feature Resolution
- Colorization
- Super-resolution
- Denoising Survey
- Super-resolution Survey

Image Denoising

Denoising

Noisy Images

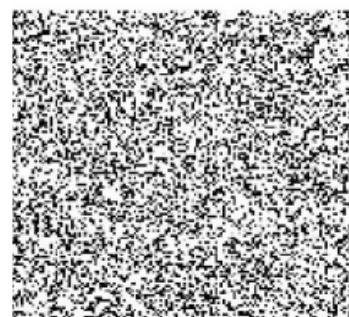
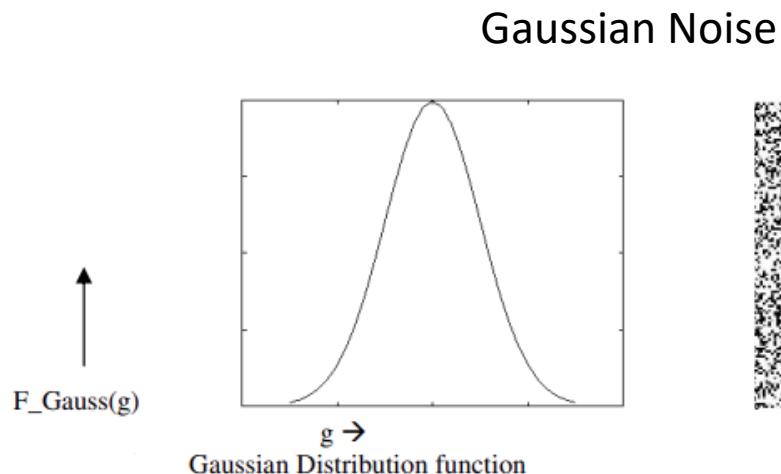


Denoised Images

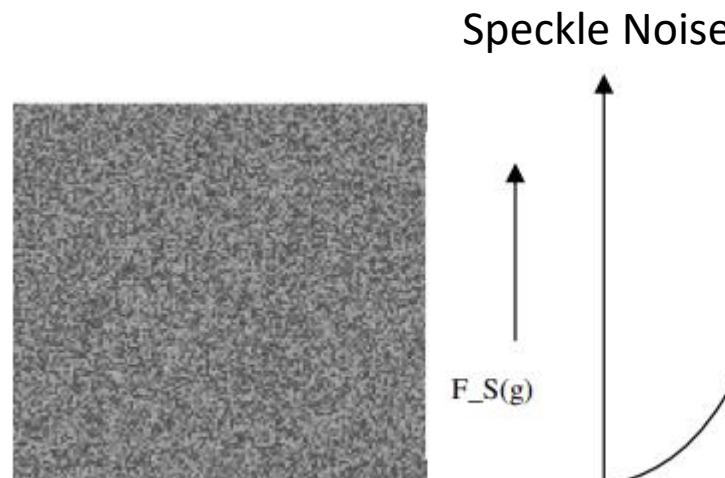
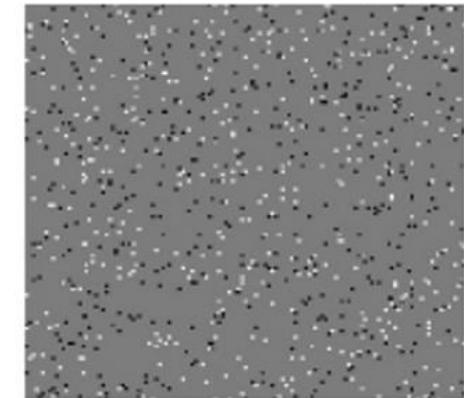
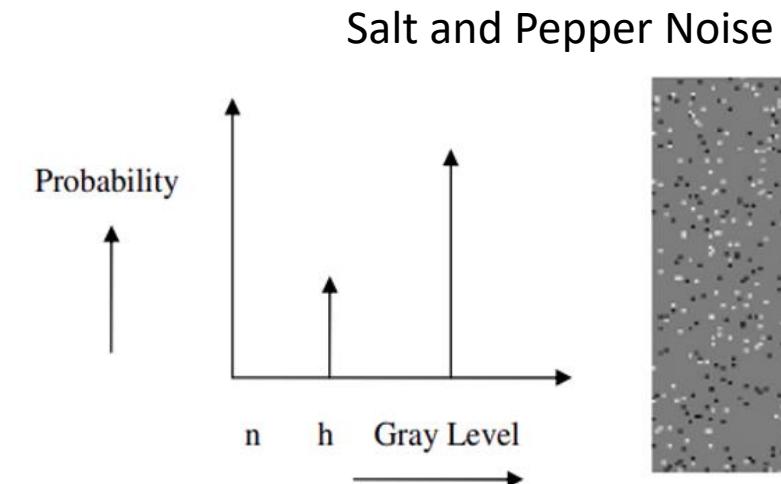


Image Denoising

Denoising



Gaussian noise



<https://analyticsindiamag.com/a-guide-to-different-types-of-noises-and-image-denoising-methods/>

Image Denoising

Traditional Image Denoising Technique

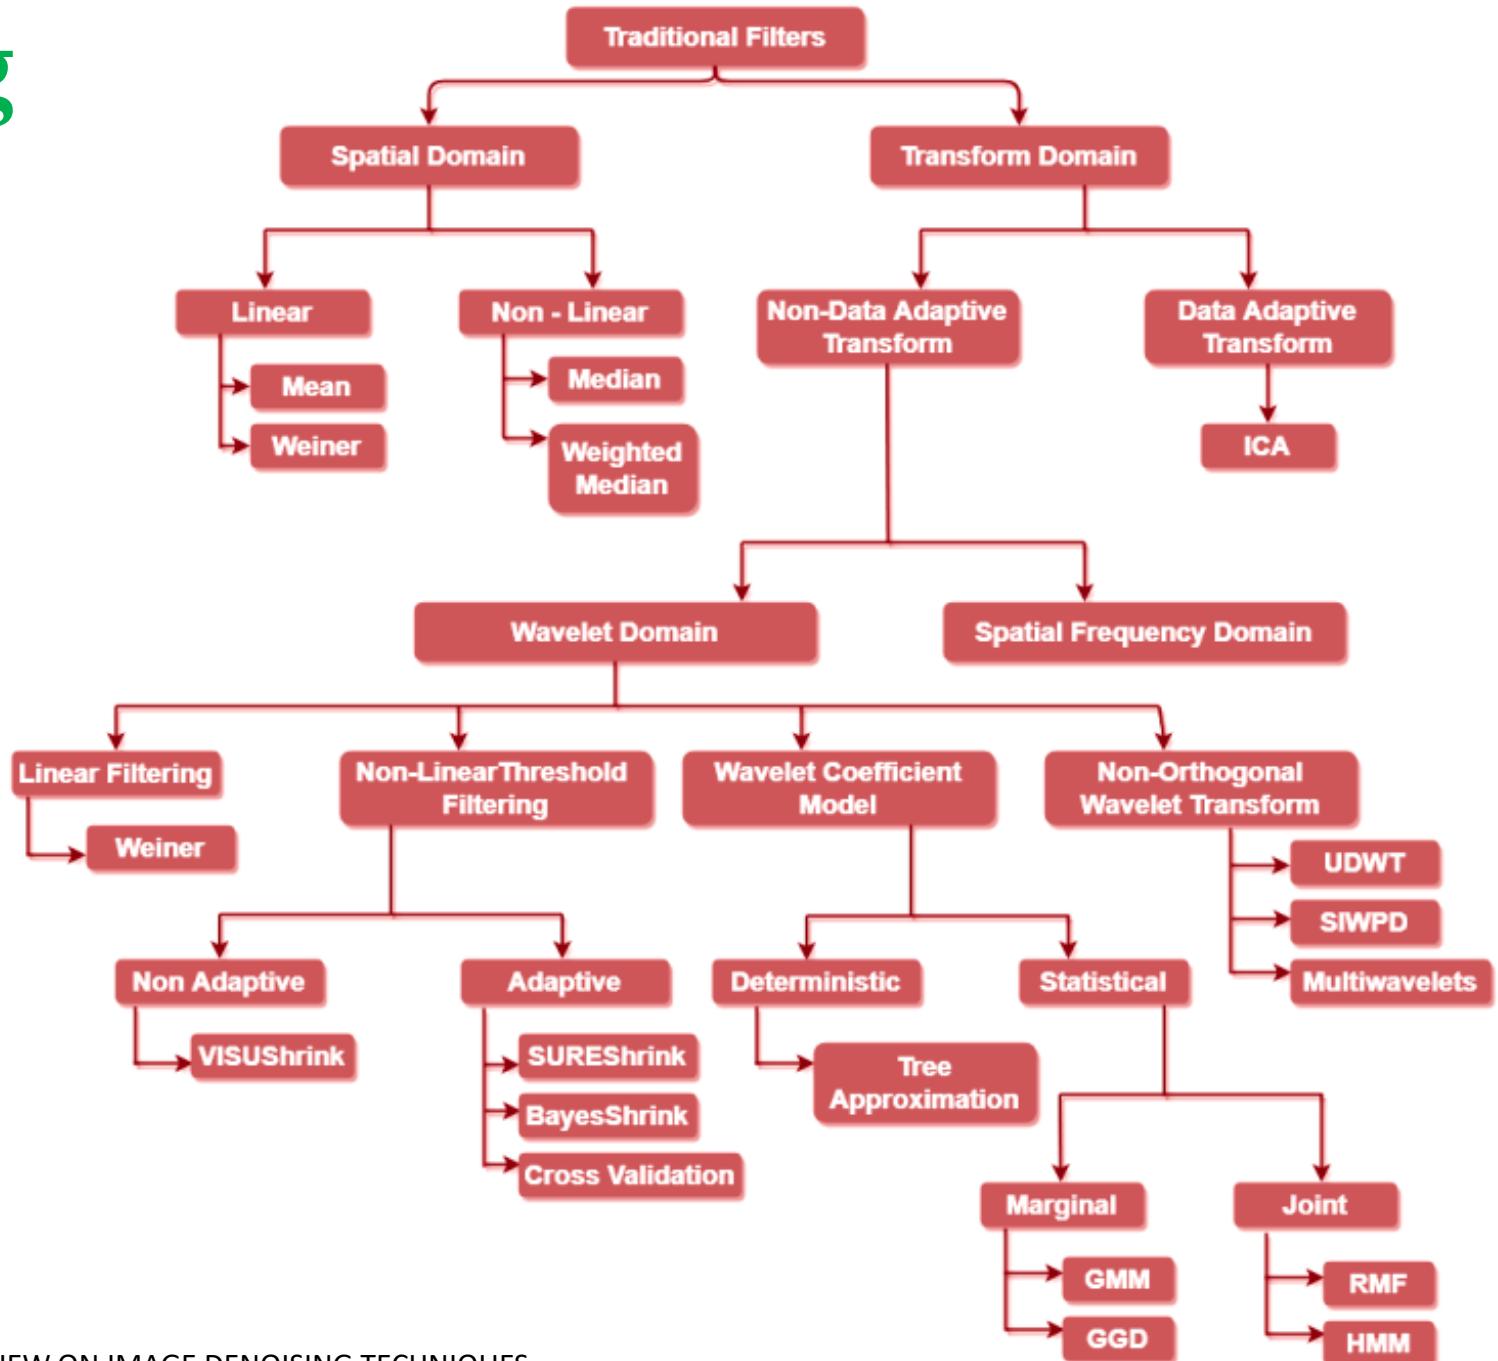
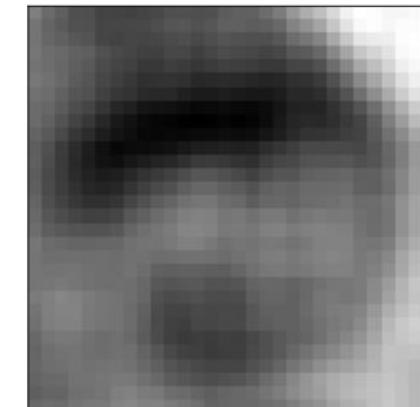
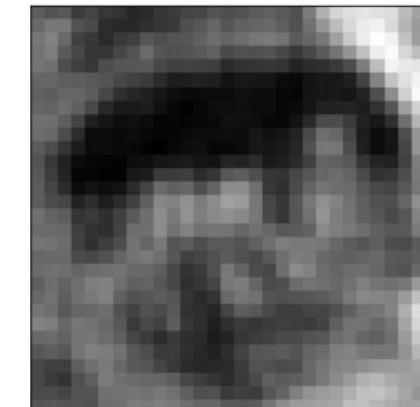
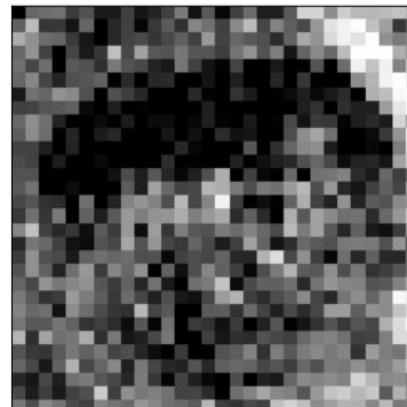
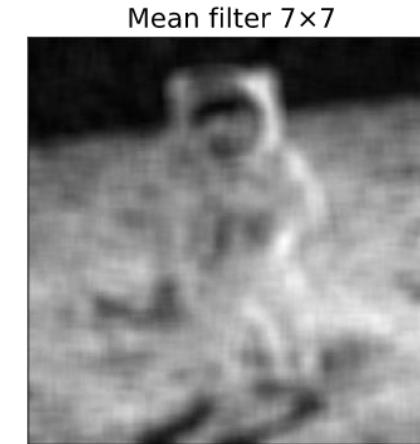
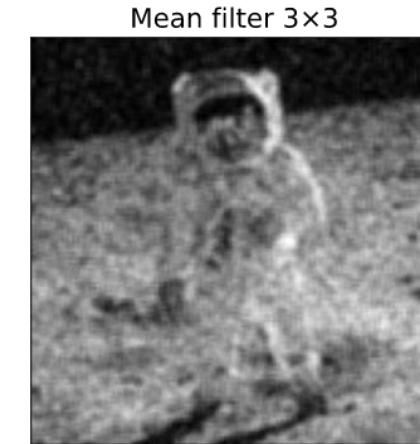
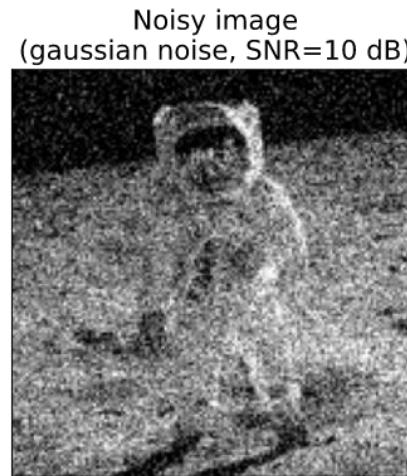


Image Denoising

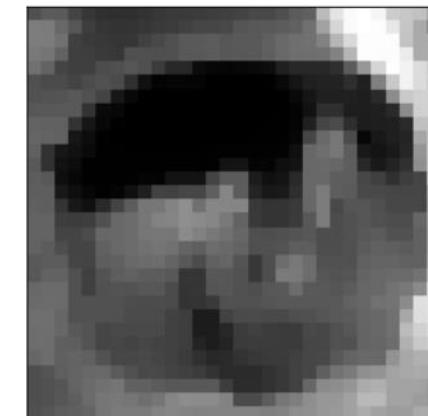
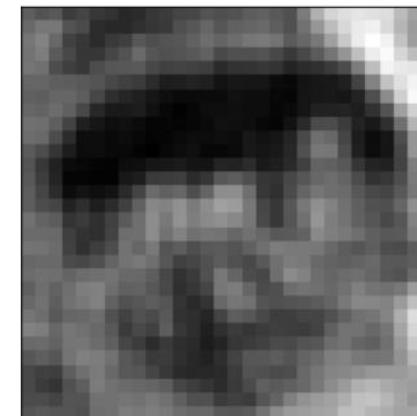
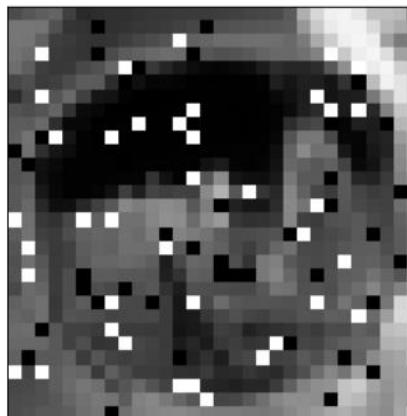
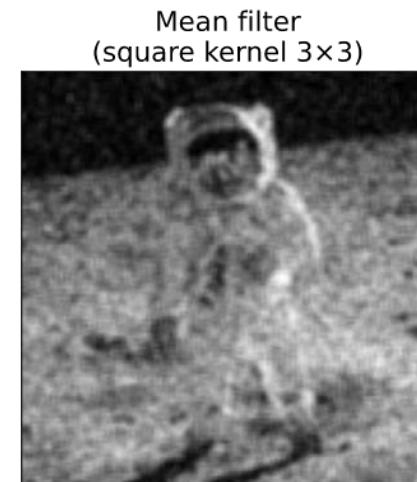
❖ Traditional Image Denoising Technique



Mean filter

Image Denoising

❖ Traditional Image Denoising Technique



Median filter

Image Denoising

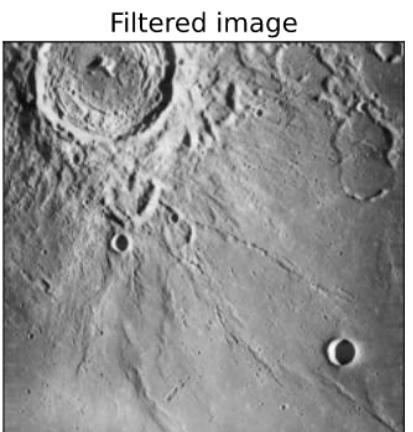
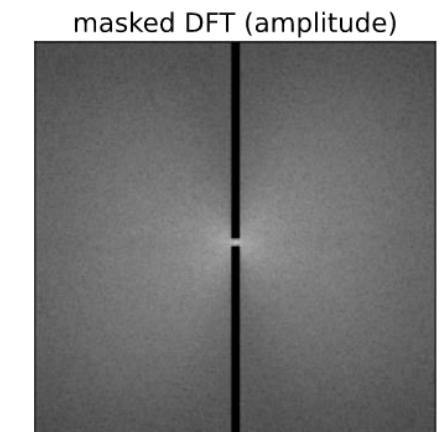
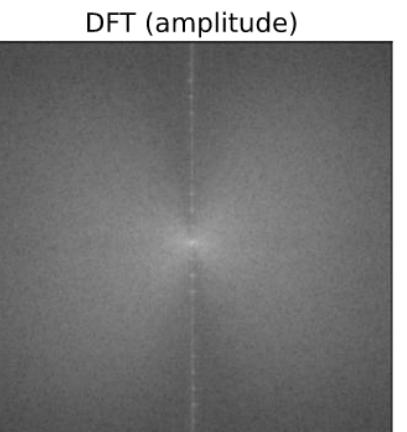
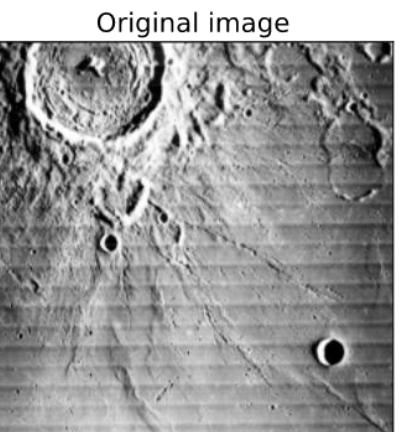
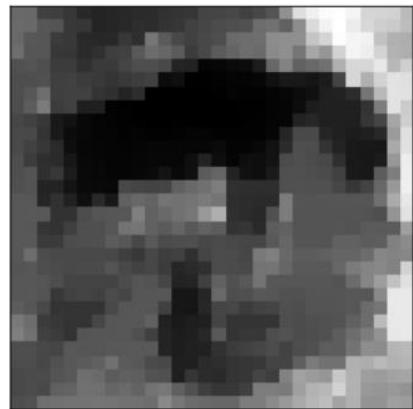
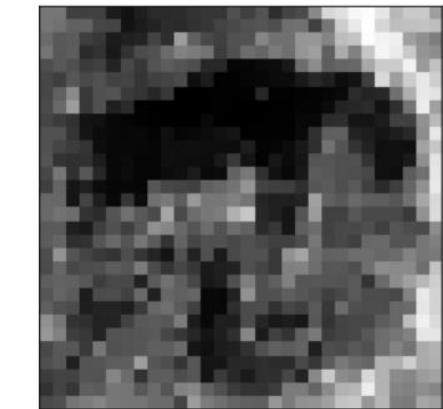
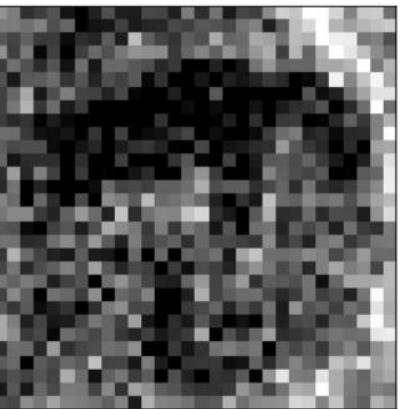
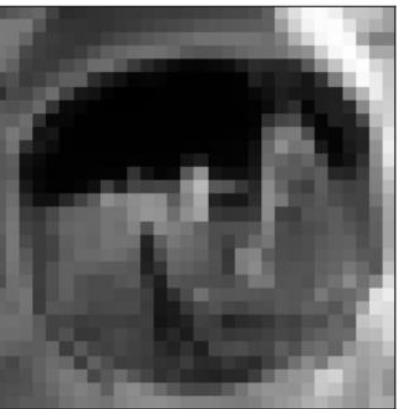
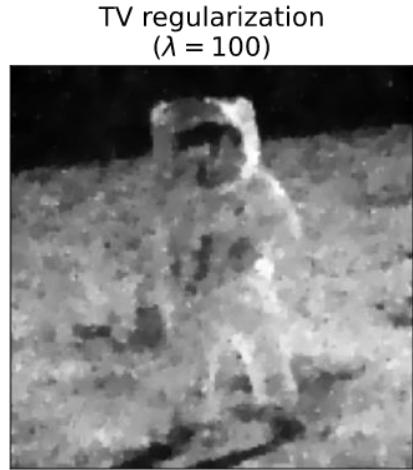
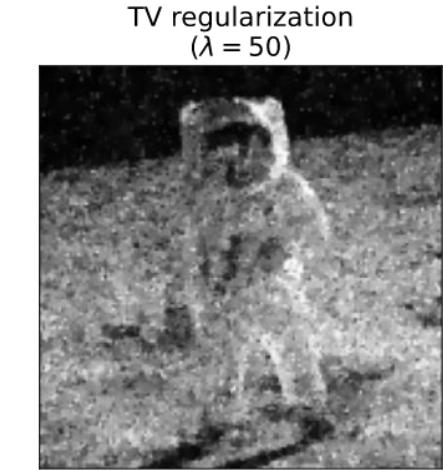


Image Denoising

❖ Based Deep-Learning Method

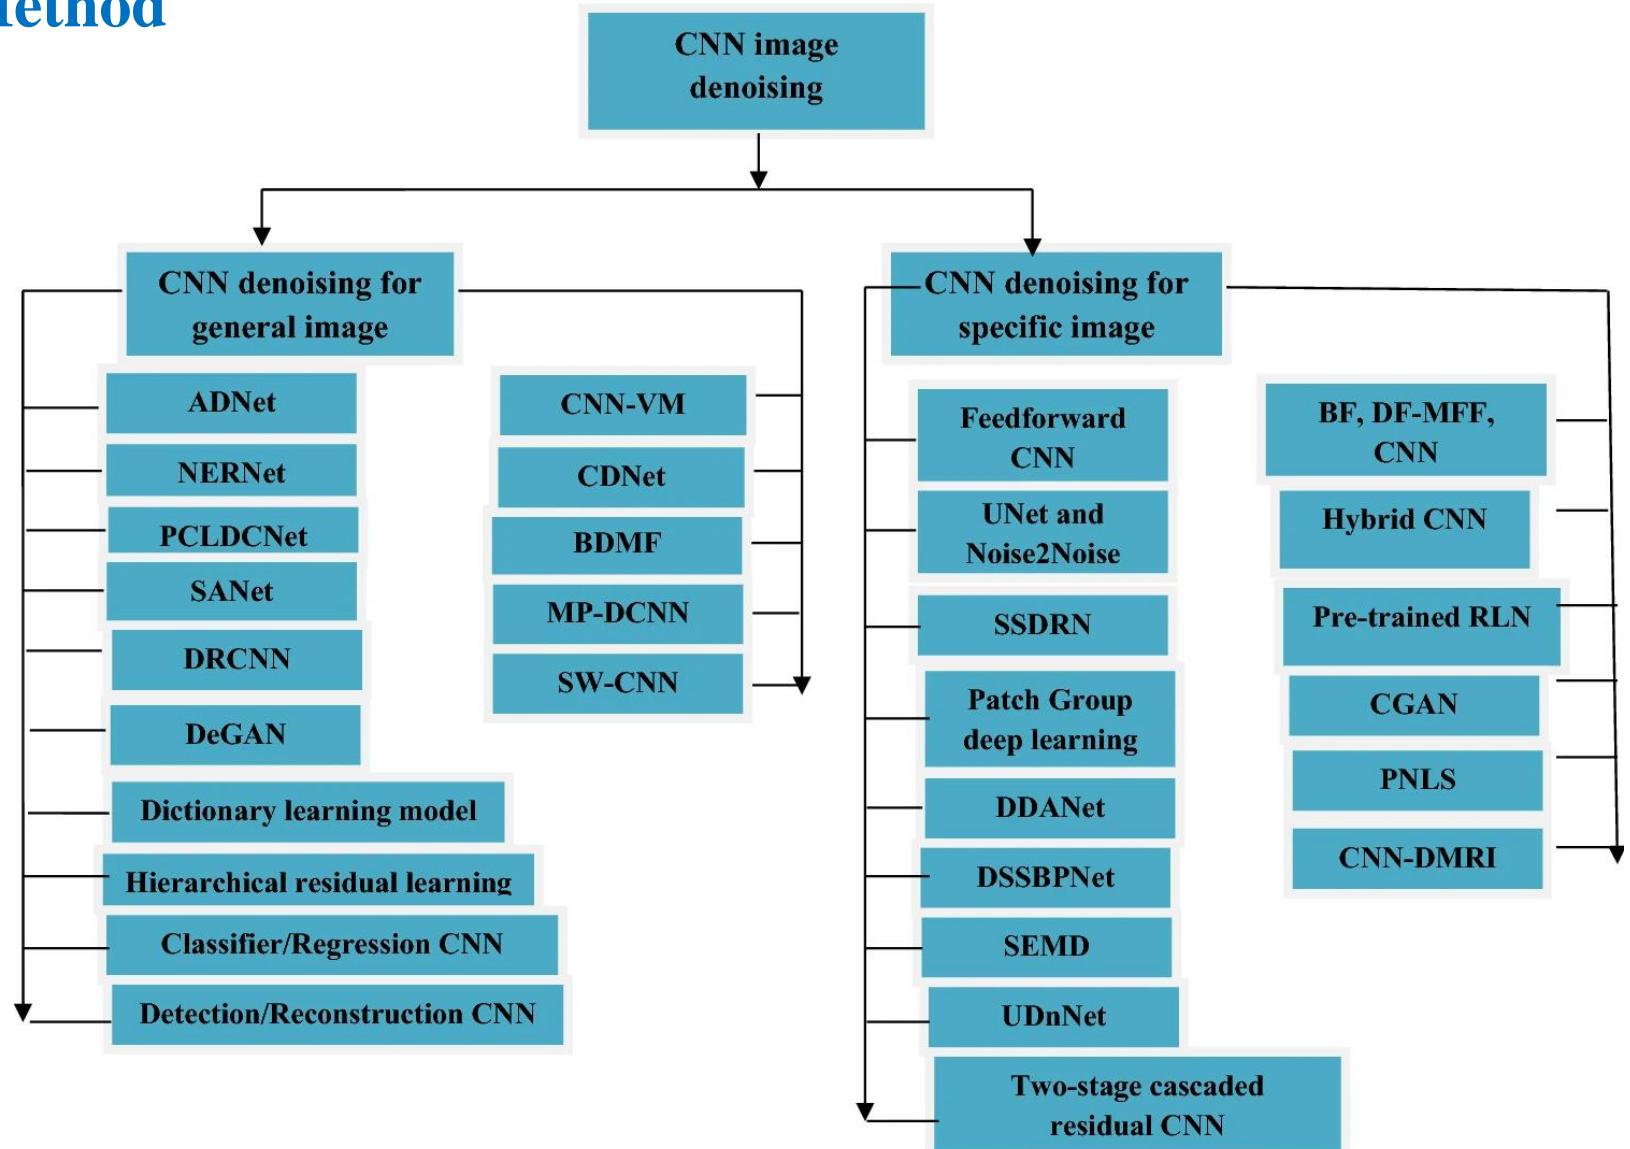
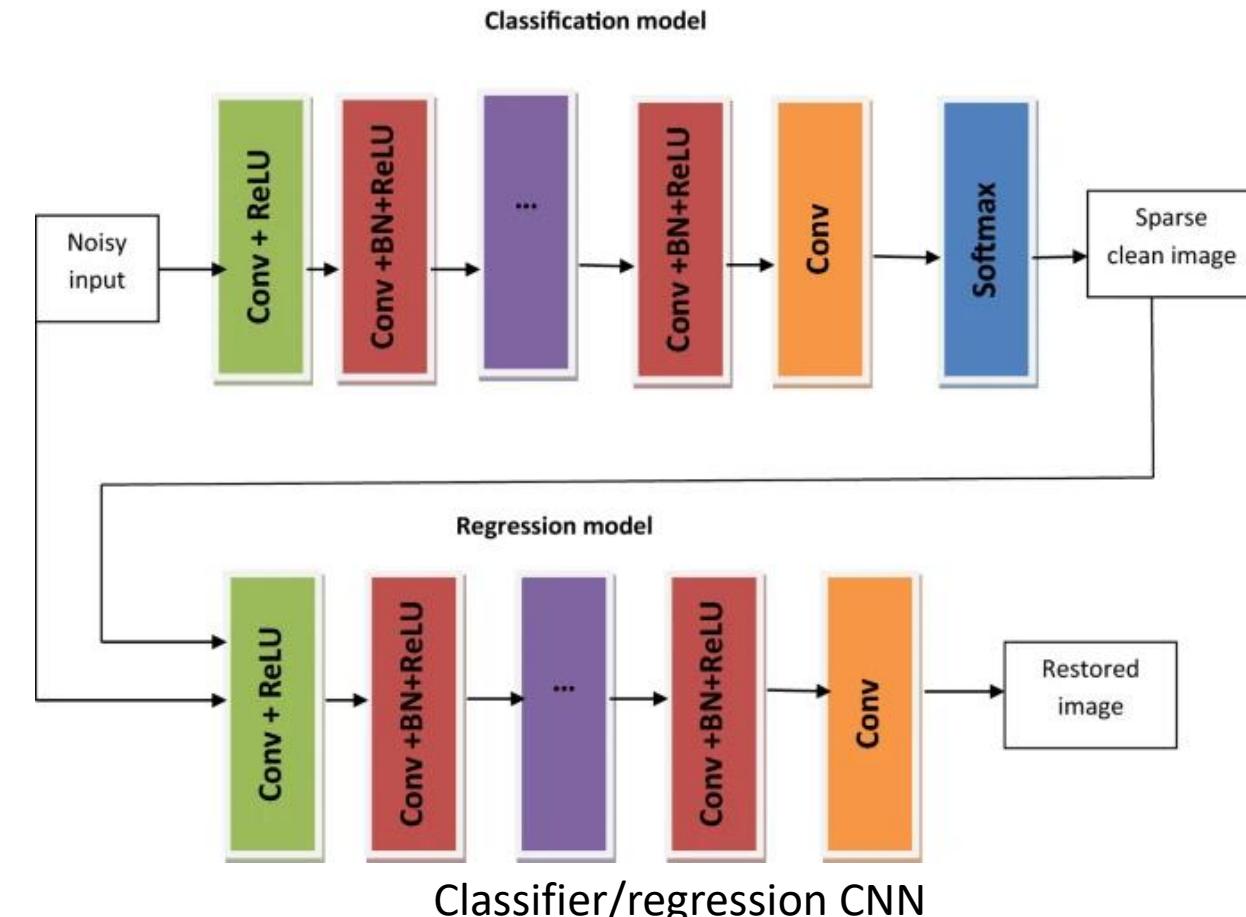
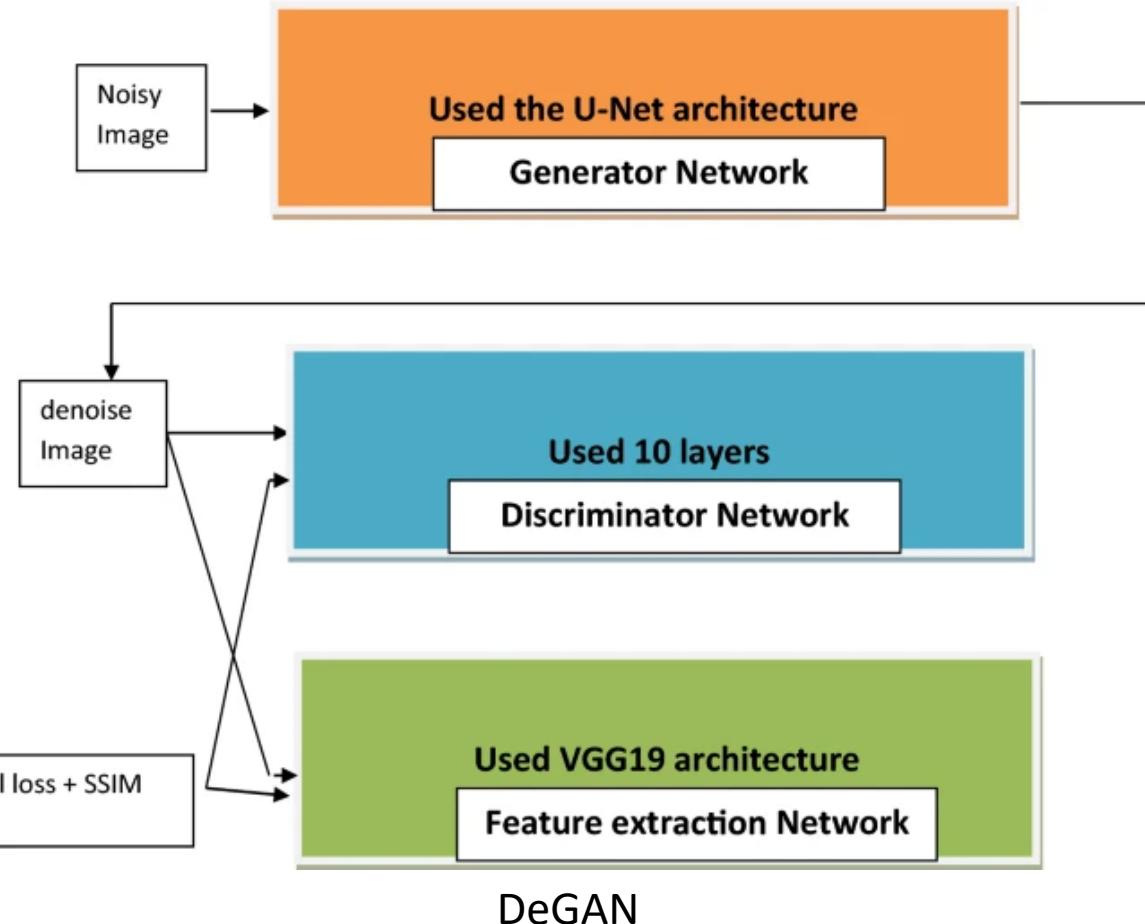


Image Denoising

❖ Based Deep-Learning Method



Outline

- Segmentation (Unet Using Interpolation)
- Alternatives to Increase Feature Resolution
- Colorization
- Super-resolution
- Denoising Survey
- Super-resolution Survey

Image Super-Resolution

<https://towardsdatascience.com/deep-learning-based-super-resolution-without-using-a-gan-11c9bb5b6cd5>

❖ Image Super-Resolution



Low-resolution



High-resolution



Low-resolution



High-resolution



Image Super-Resolution

❖ Image Super-Resolution

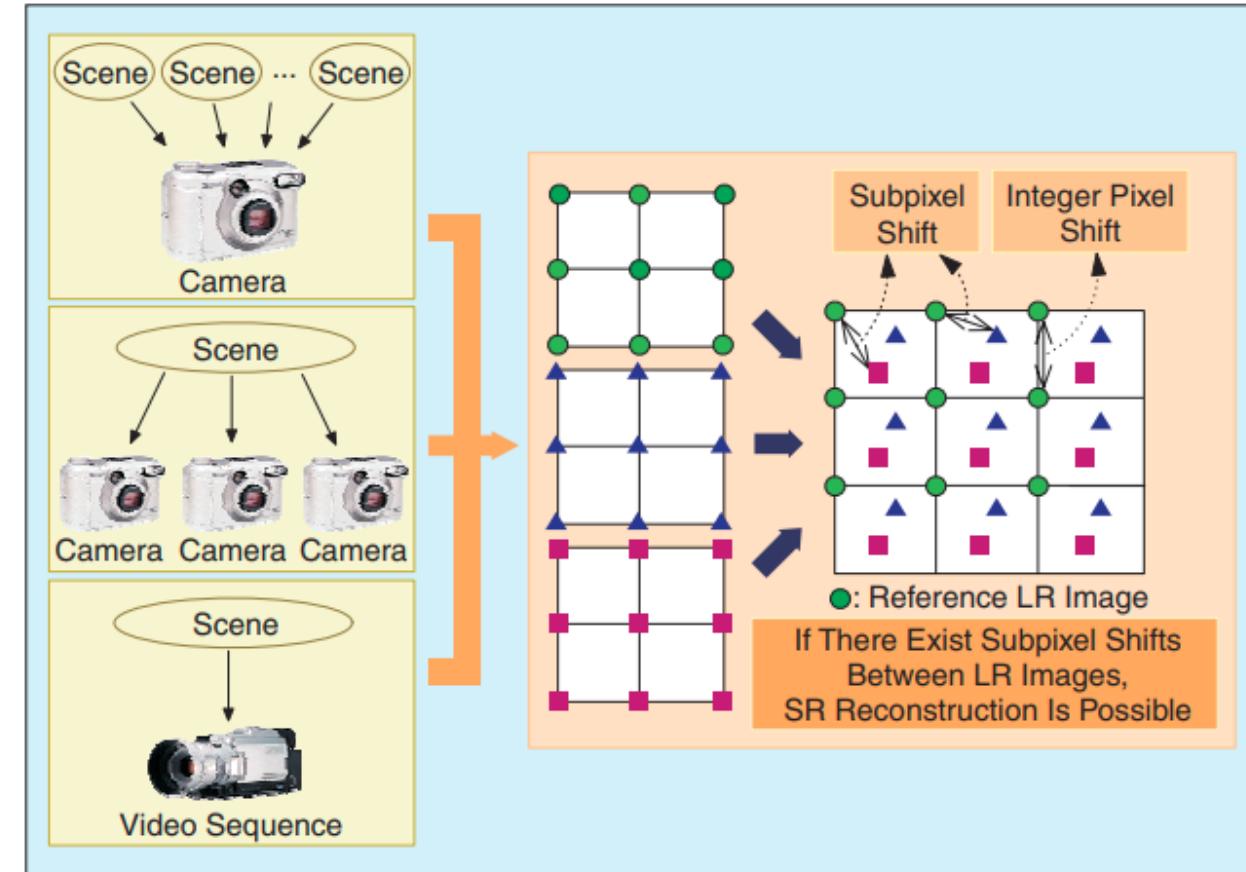
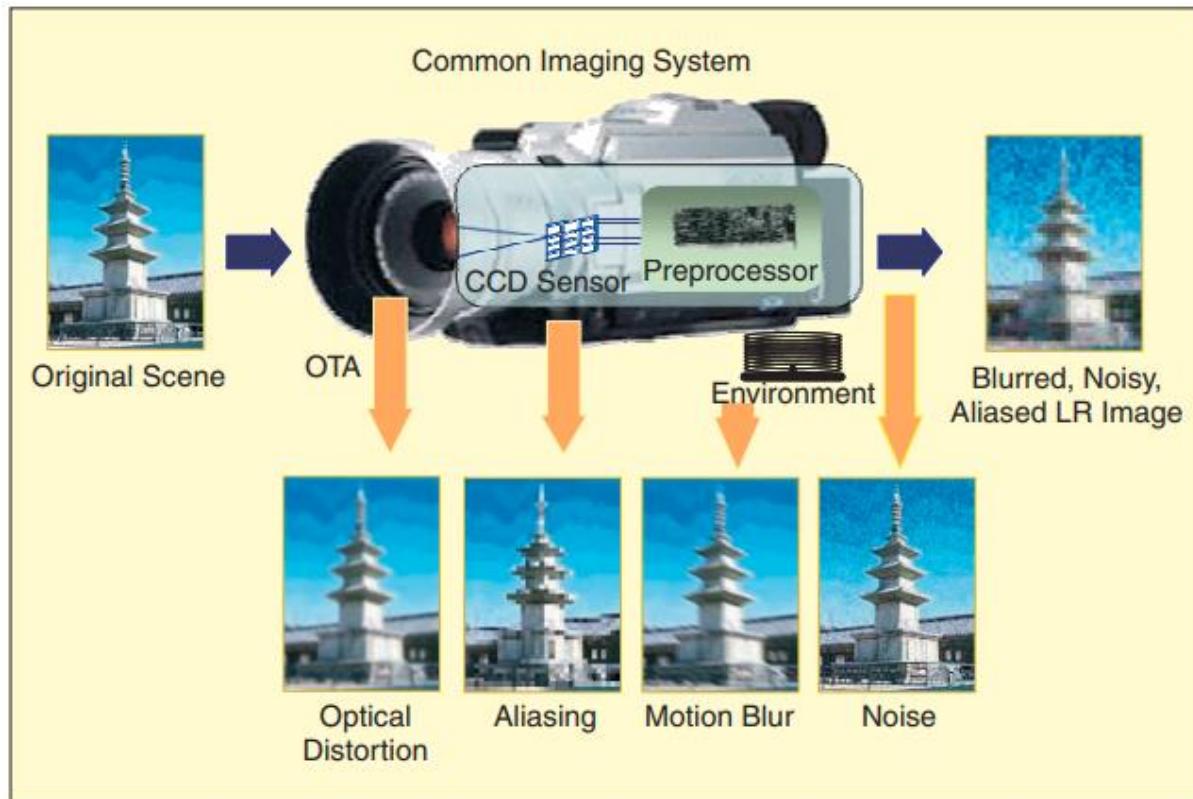
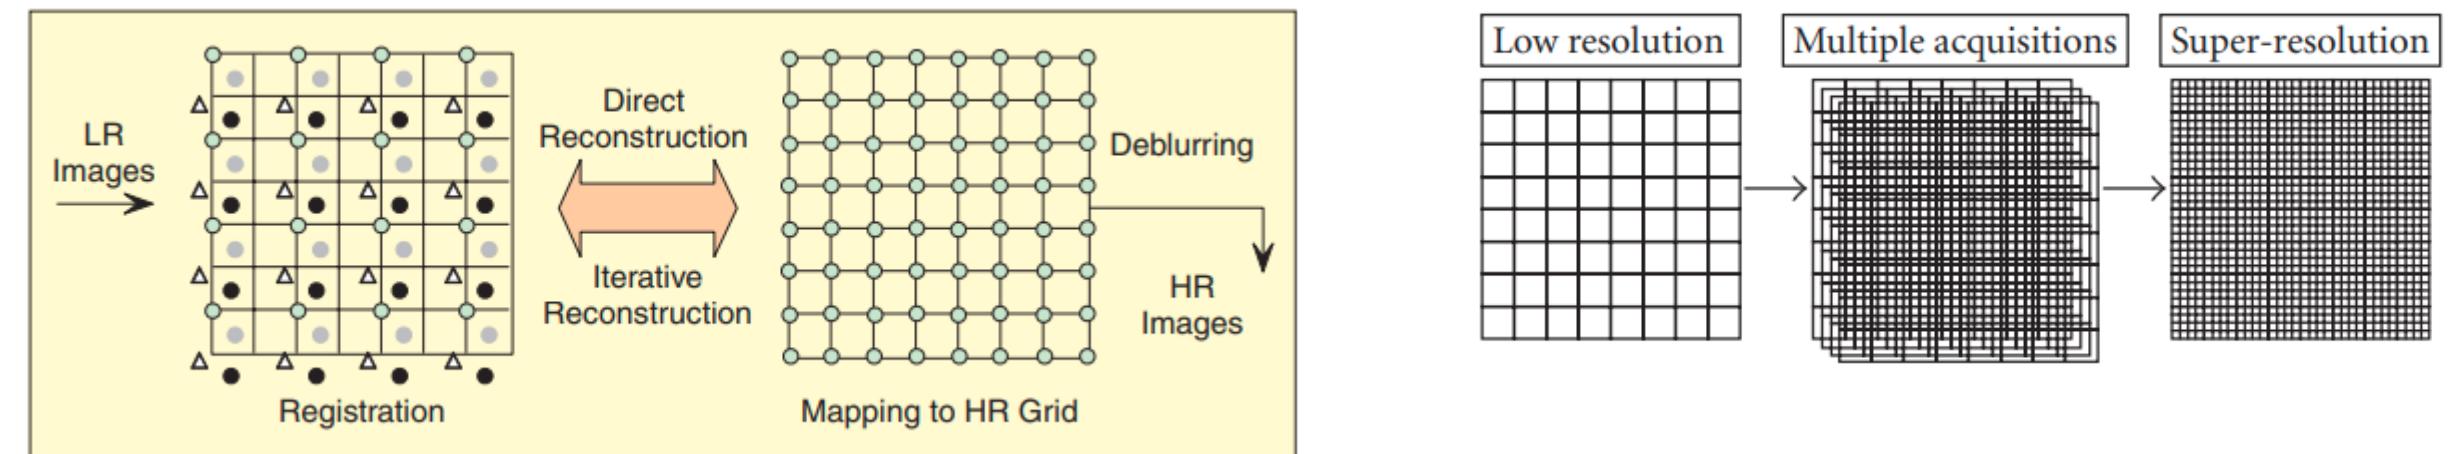
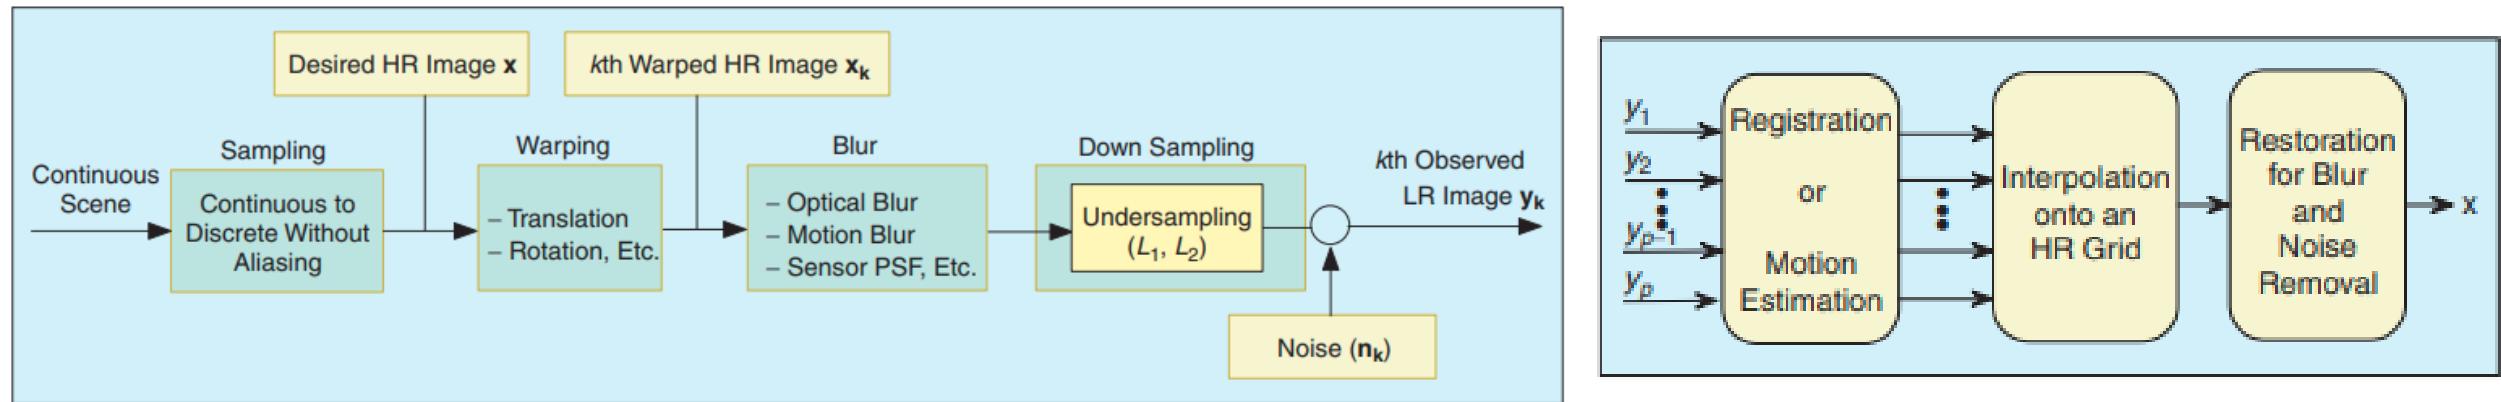


Image Super-Resolution

❖ Traditional Method



Park, S.C., Park, M.K., Kang, M.G.: Super-resolution image reconstruction: a technical overview. IEEE Signal Process. Mag. 20(3), 21–36 (2003)

Kennedy, John & Israel, Ora & Frenkel, Alex & bar-shalom, Rachel & Azhari, Haim. (2007)

Improved Image Fusion in PET/CT Using Hybrid Image Reconstruction and Super-Resolution. International journal of biomedical imaging. 2007

Image Super-Resolution

❖ Based Deep Learning Method

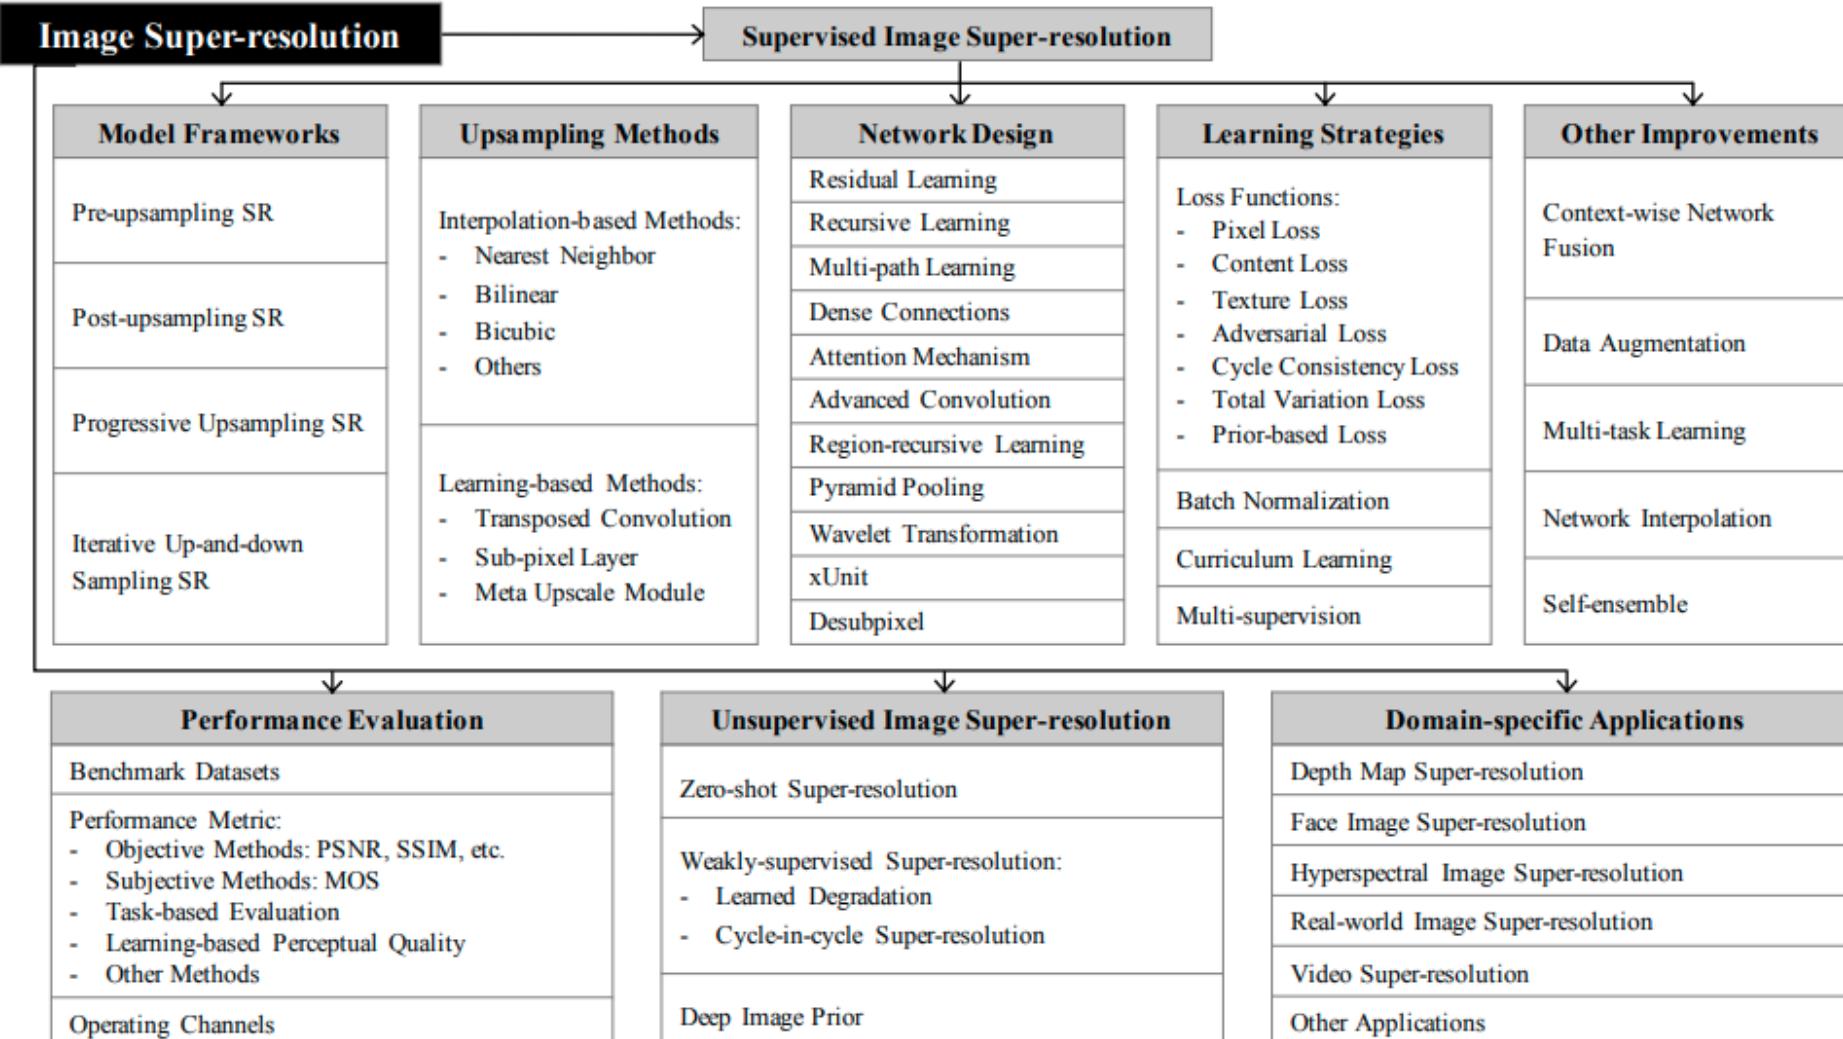
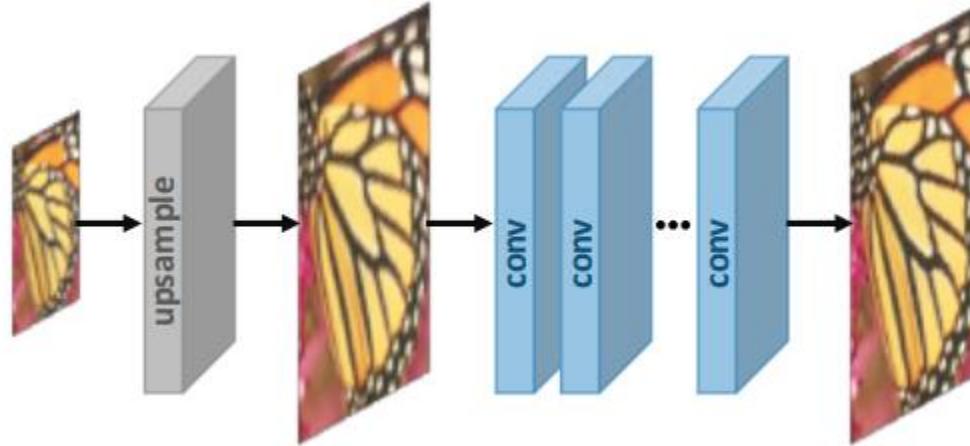


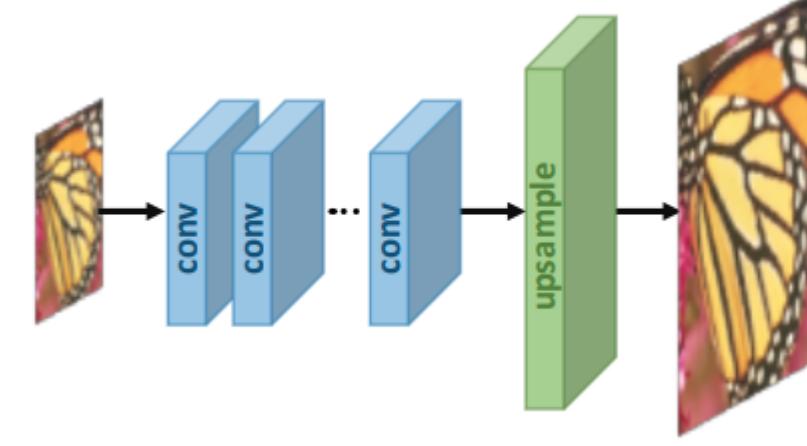
Image Super-Resolution

❖ Model Frameworks

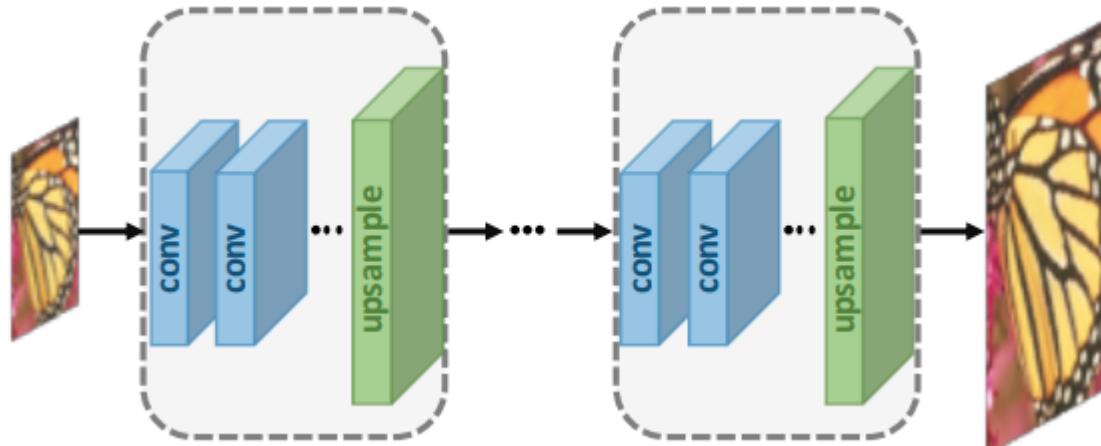
Pre-upsampling SR



Post-upsampling SR



Progressive upsampling SR



Iterative up-and-down Sampling SR

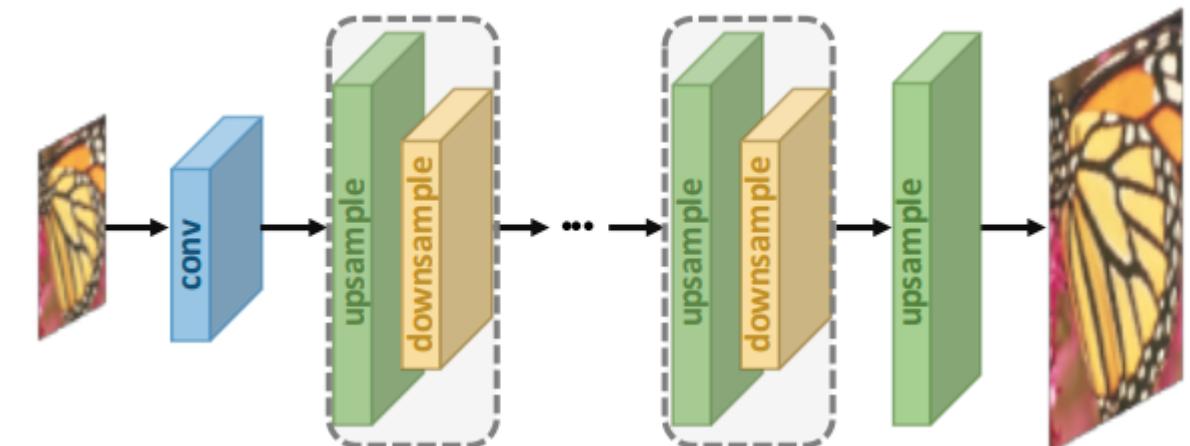
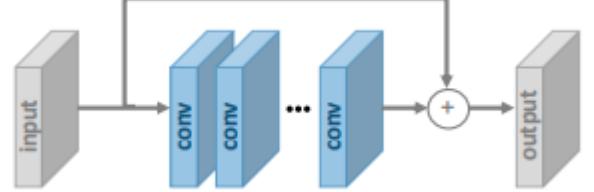


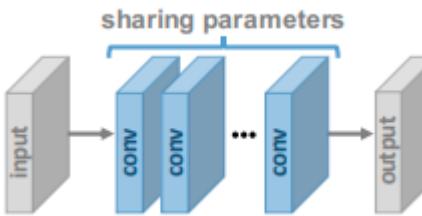
Image Super-Resolution

❖ Network Design

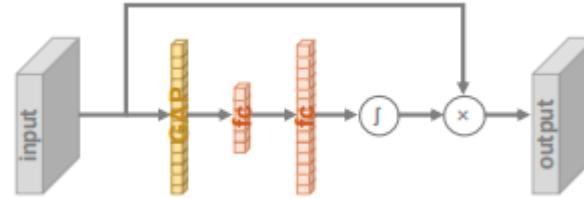
Residual Learning



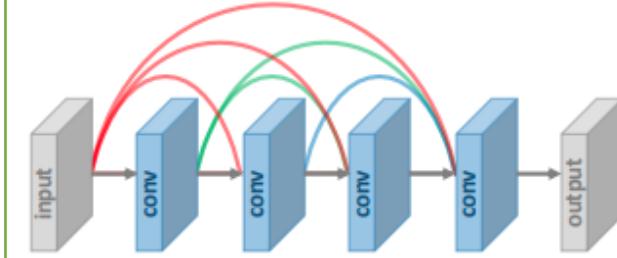
Recursive learning



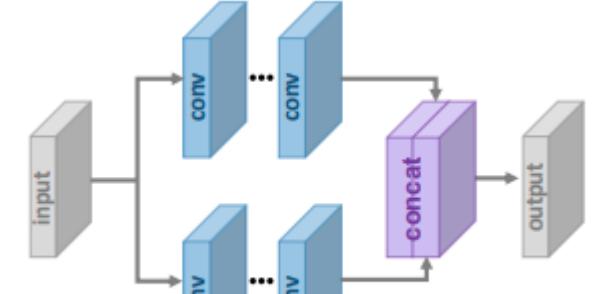
Channel attention



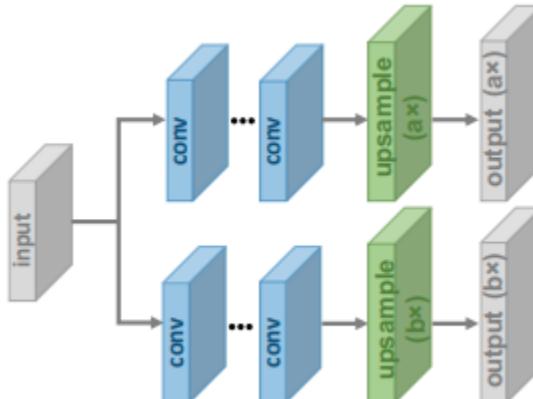
Dense connections



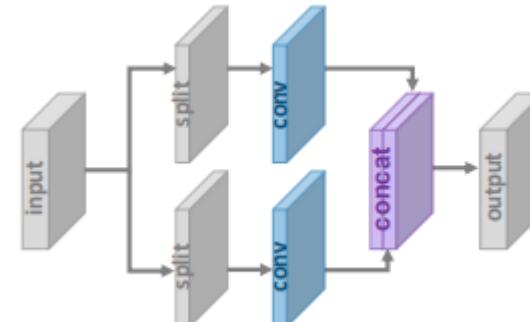
Local multi-path learning



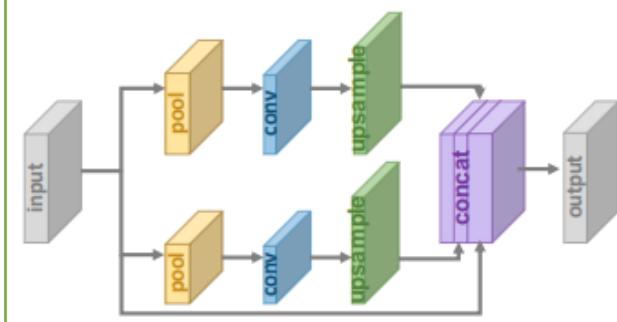
Scale-specific multi-path learning



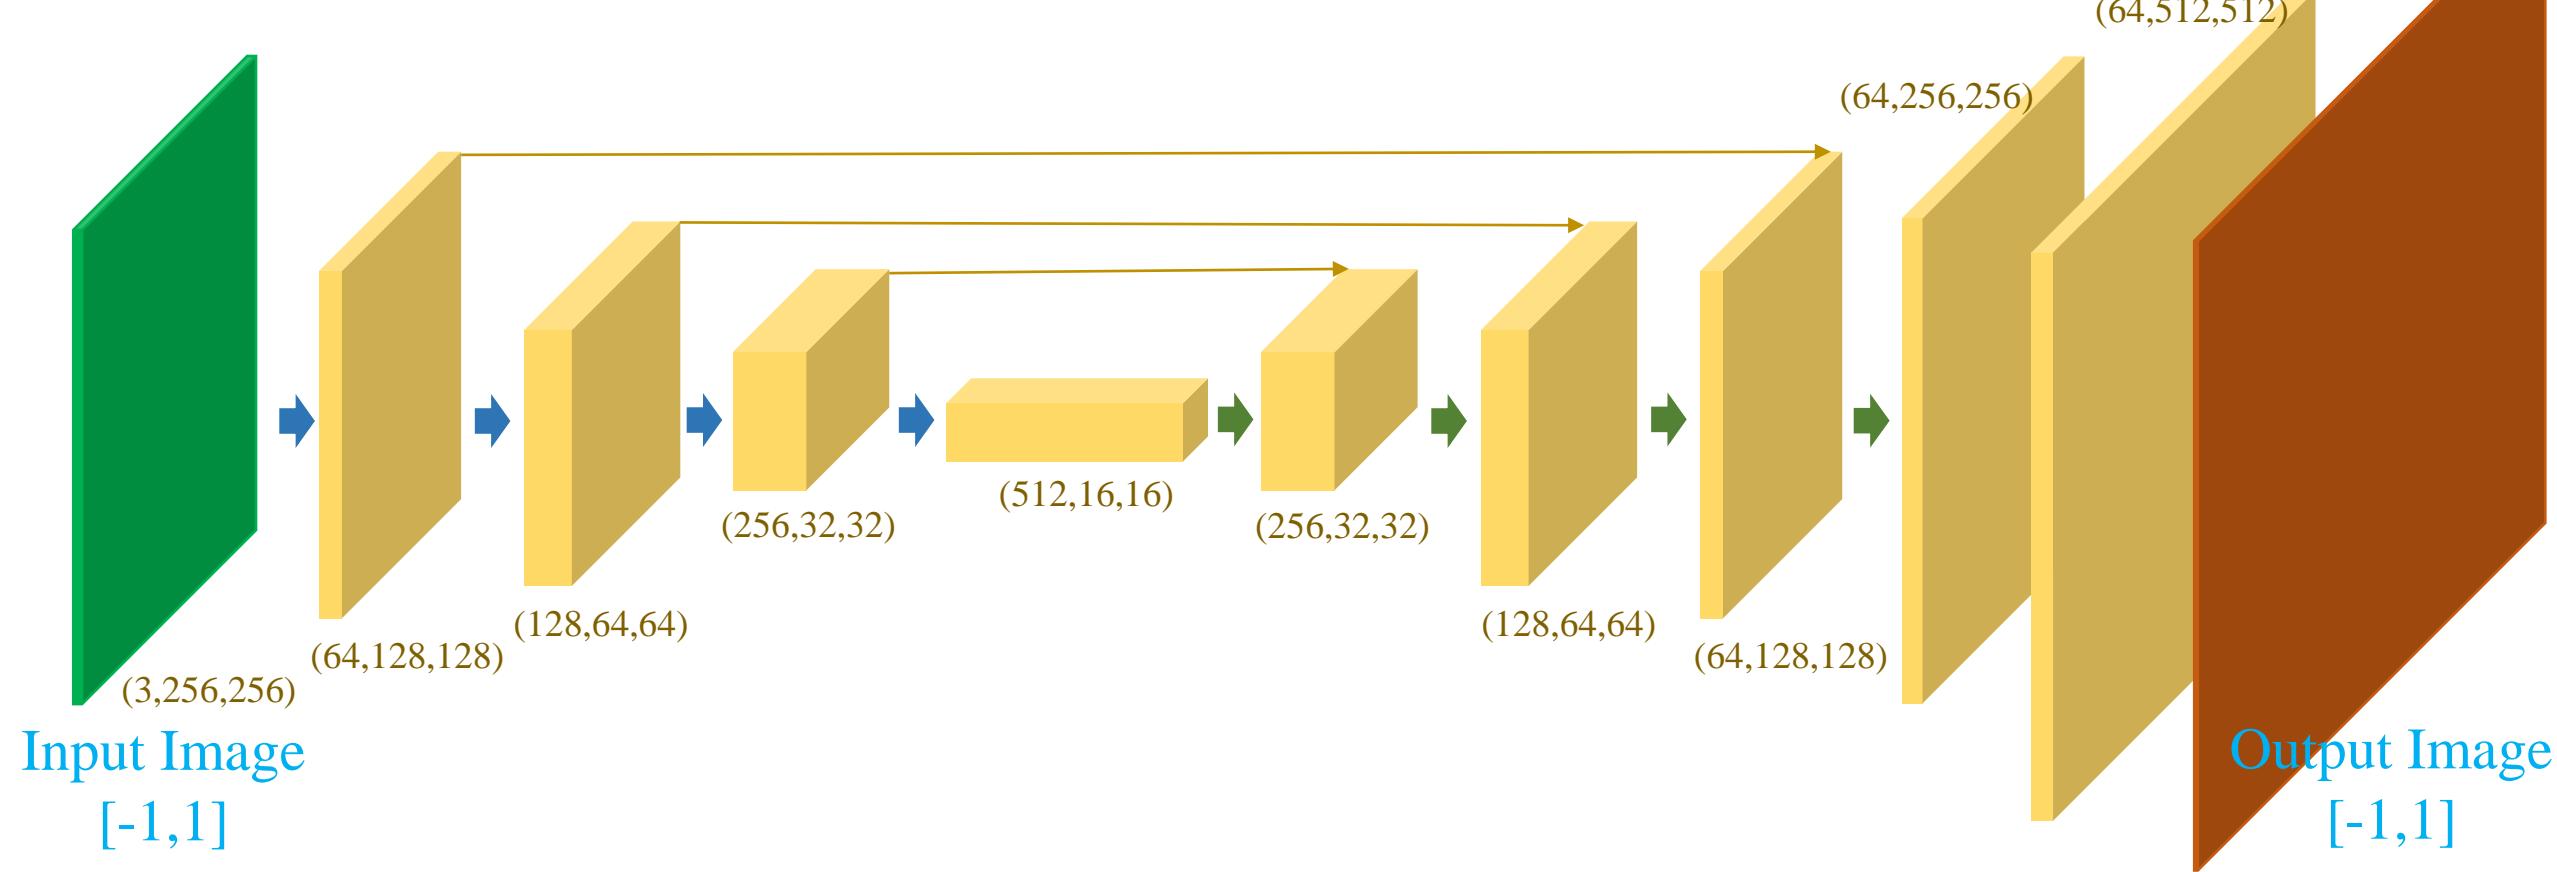
Group convolution



Pyramid pooling



Super-resolution



Blue Block:
 (3×3) Convolution
padding = 'same'
stride = 1 + ReLU
+ Batch Norm
+ (2×2) max pooling

Green Block:
 4×4 ConvTranspose
padding = 'same'
stride = 2 + ReLU
+ Batch Norm

Red Block:
 4×4 ConvTranspose
padding = 'same'
stride = 2 + Tanh

